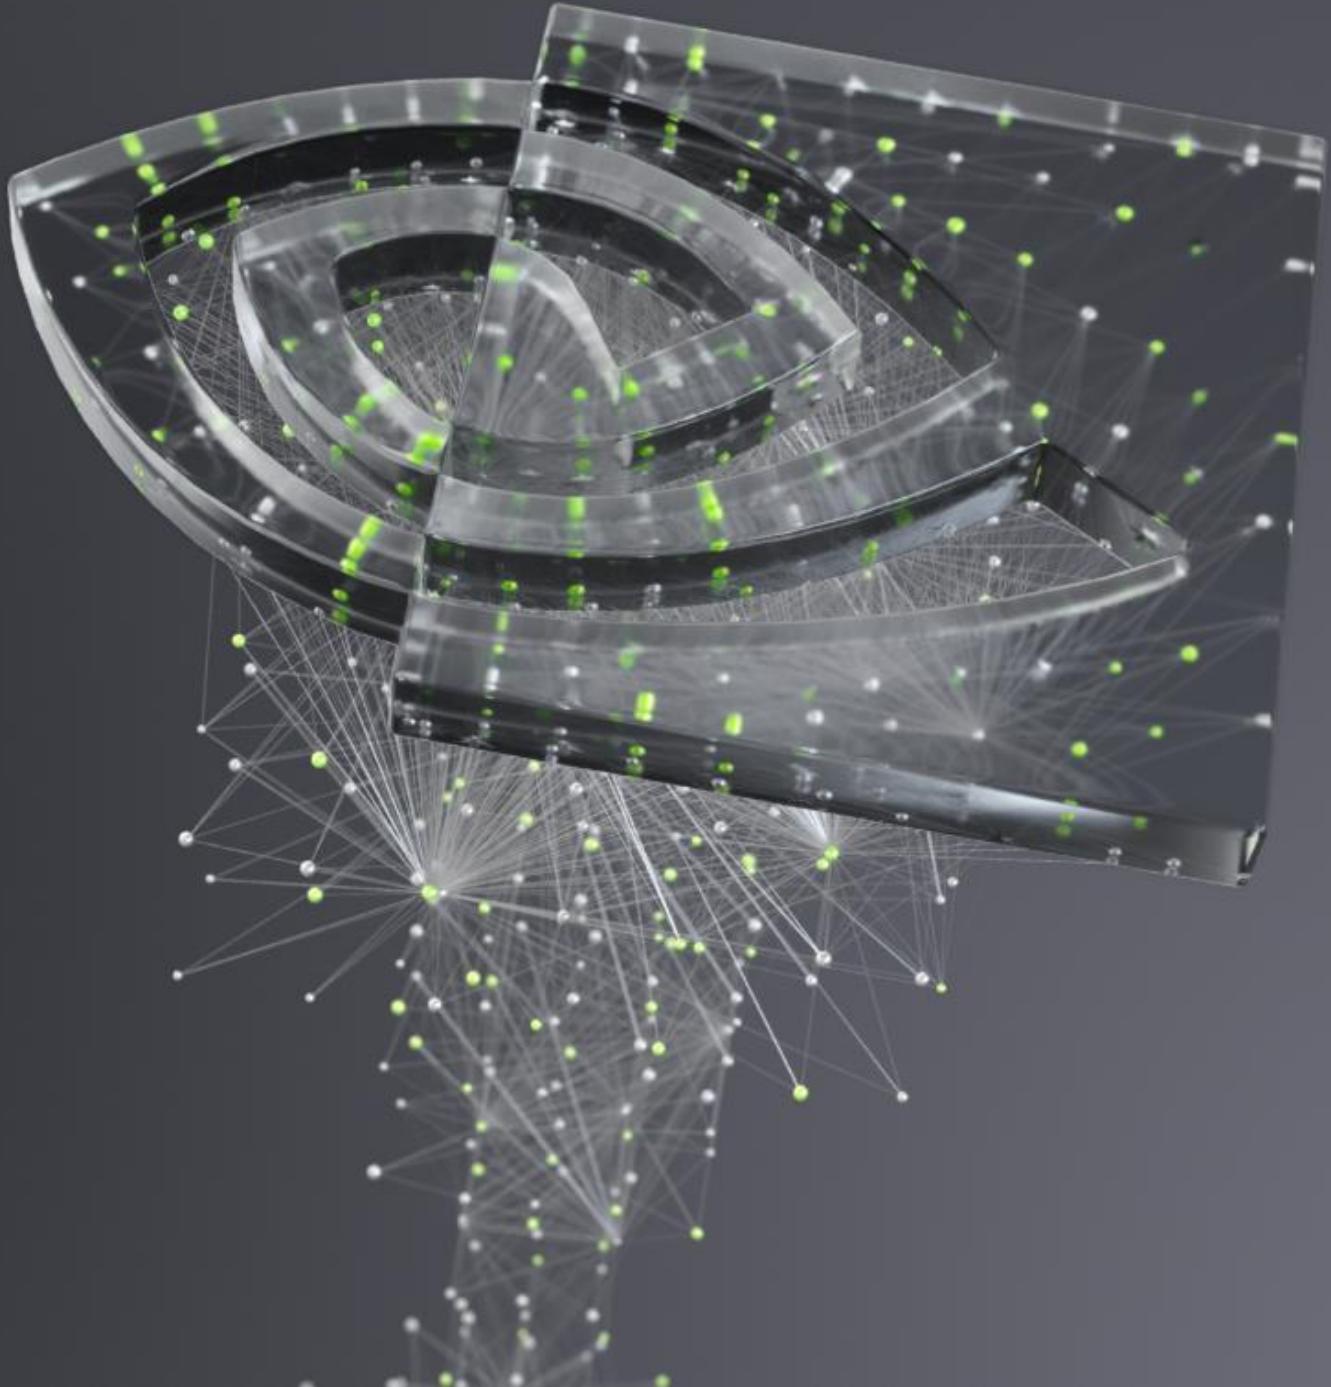




JETSON PLATFORM

Jetson SA: Alan Zhang

April 24, 2020



AGENDA

Jetson product and marketing

Why we need cross compile

How to setup the cross-compile environment

1. Config directly via the environment variable.
 2. Config in Makefile.
 3. Build your simulation.
-

Compile on X86 host and execute on Jetson

1. Cross compile Linux kernel.
2. Cross compile CUDA project.
3. Cross compile Multi Media API based application.
4. Cross compile Deepstream.
5. Cross compile Caffe
6. Cross compile OpenCV.

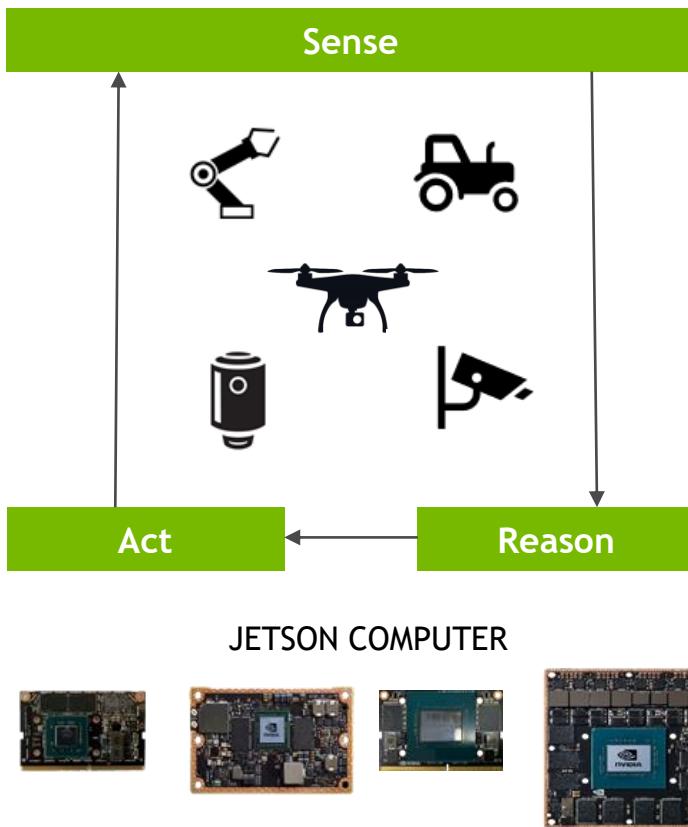
Where can I get the prebuild version and how to build on jetson?

NVIDIA JETSON

Software-Defined AI Platform

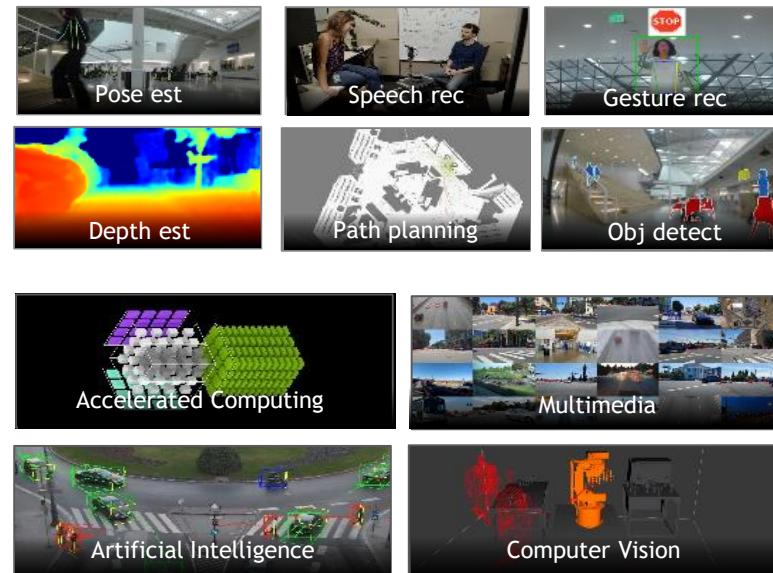
AI at the Edge

Sensor Fusion & Compute Performance



SOFTWARE DEFINED

SDK, Design Tools, Libs, GEMs



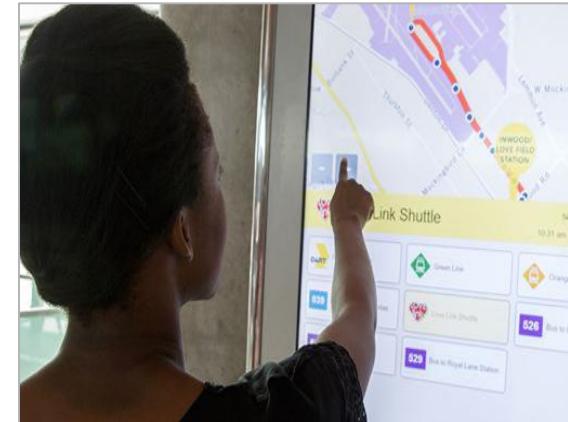
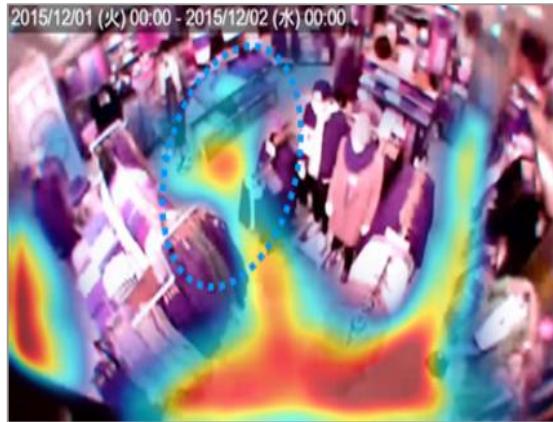
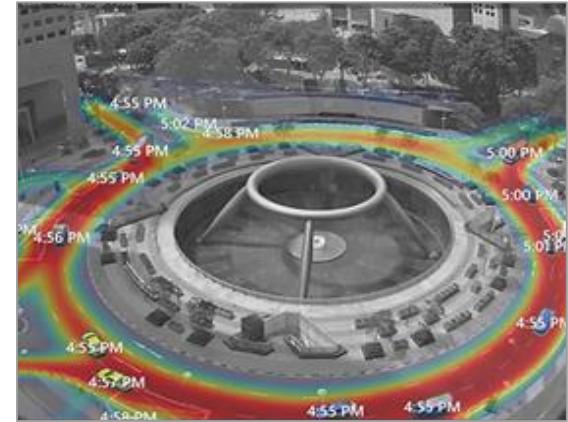
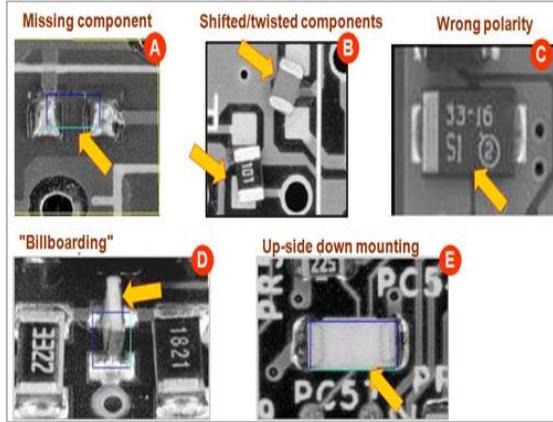
Jetpack SDK · CUDA · TensorRT · TensorFlow · ONNX · ROS

ECOSYSTEM

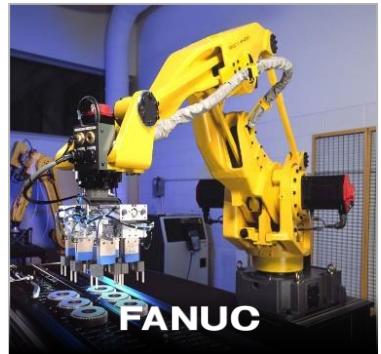
Expertise, Time to Market



INTELLIGENT VIDEO ANALYTICS FOR EFFICIENCY AND SAFETY



JETSON ROBOTICS & LOGISTIC



Industrial



Aerospace/Defense



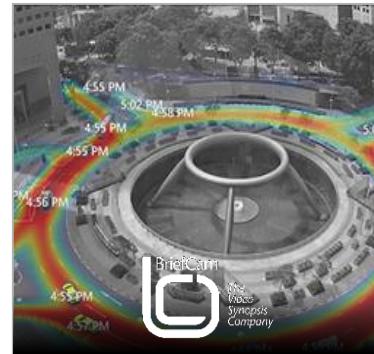
Construction



Agriculture



Healthcare



Smart City



Retail



Logistics



Delivery



Inspection



Service



Collaboration

THE JETSON FAMILY

for AI at the Edge and Autonomous System designs

JETSON NANO

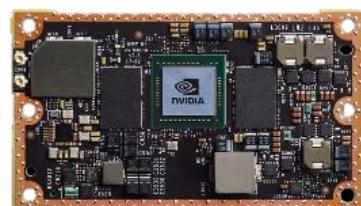
0.5 TFLOPS (FP16)



5 - 10W
45mm x 70mm

JETSON TX2 series

1.3 TFLOPS (FP16)



7.5 - 15W*
50mm x 87mm

JETSON Xavier NX

6 TFLOPS (FP16)
21 TOPS (INT8)



10 - 15W
45mm x 70mm

JETSON AGX XAVIER series

11 TFLOPS (FP16)
32 TOPS (INT8)



10 - 30W
100mm x 87mm

AI at the edge

Fully autonomous machines

Same software

Listed prices are for 1000u+ | Full specs at developer.nvidia.com/jetson

* TX2i: 10-20W



Jetson AGX Xavier Developer Kit

- [Developer Kit User Guide](#)
- [L4T Documentation](#)
- [Module Datasheet](#)
- [OEM Design Guide](#)



Jetson TX2 Developer Kit

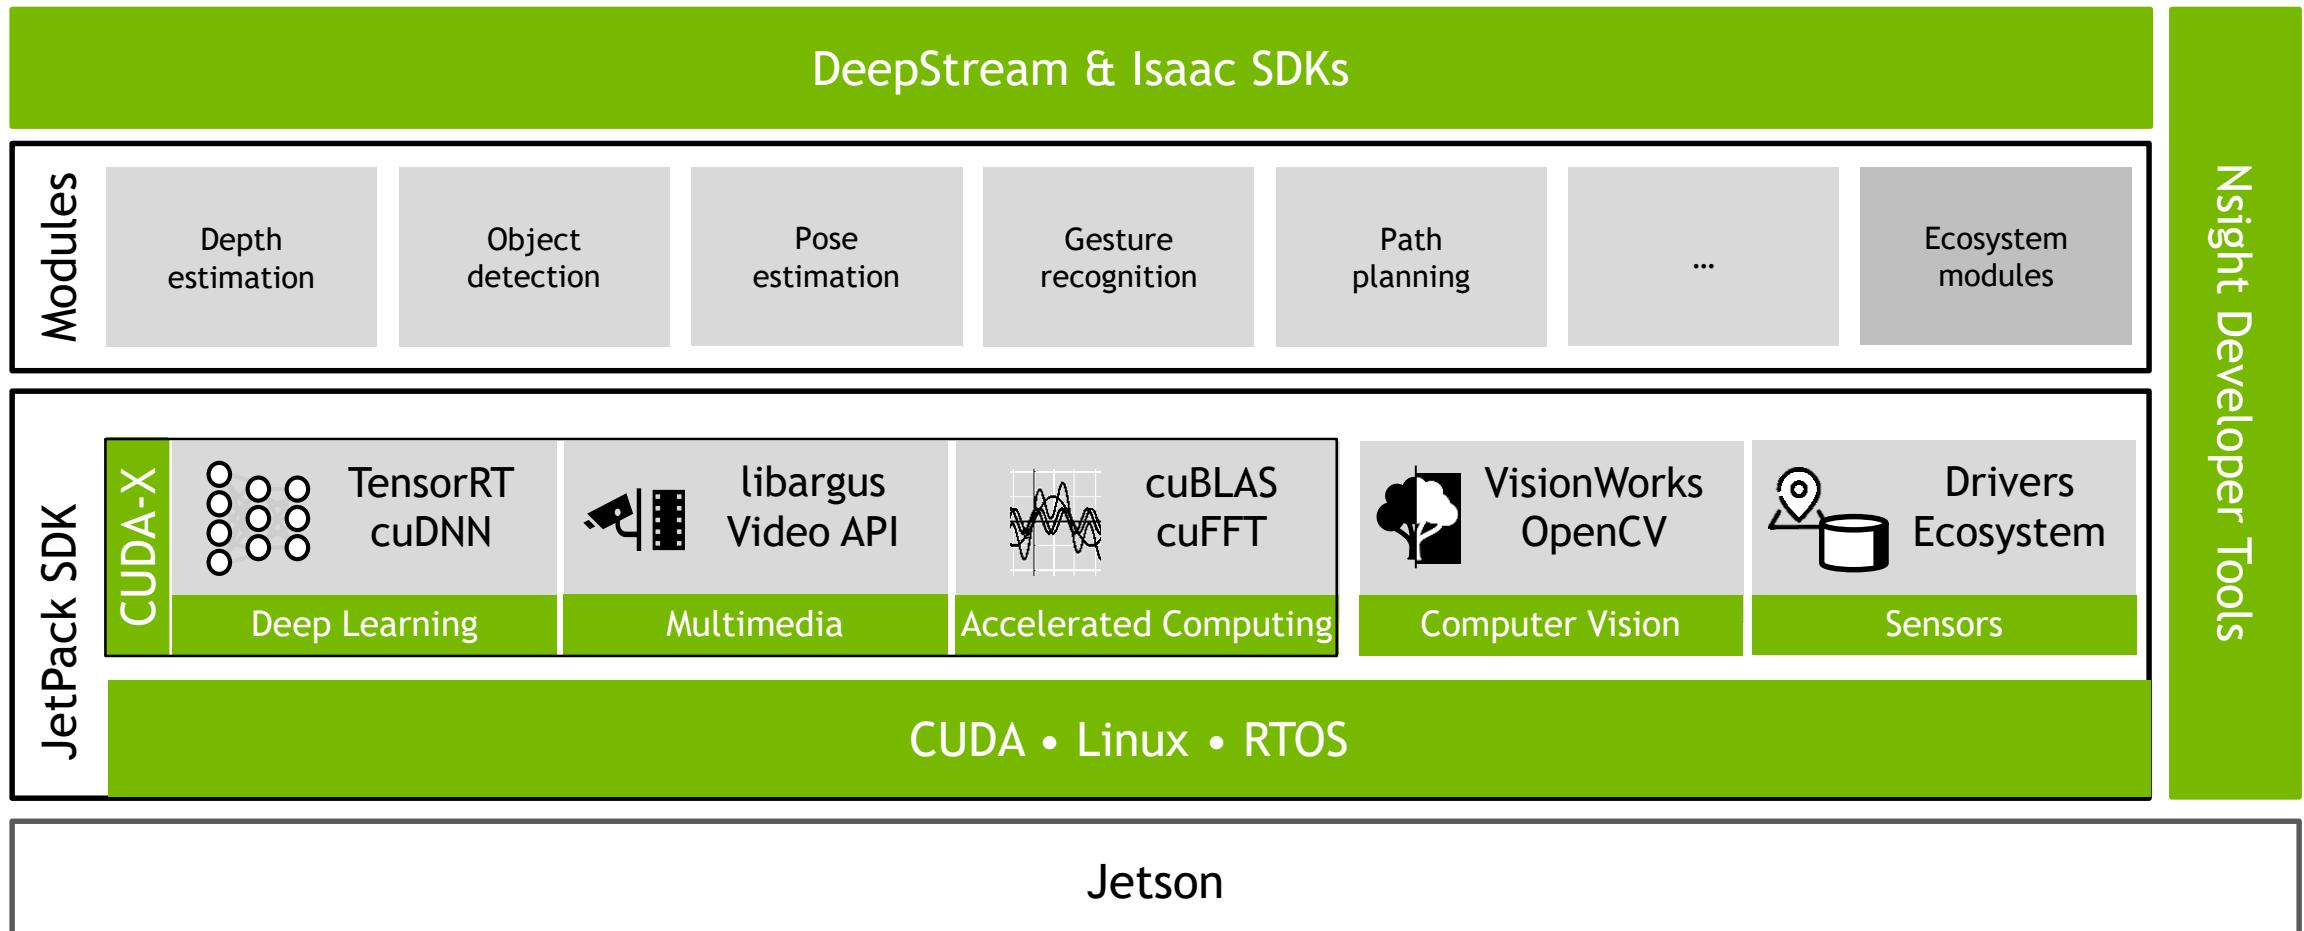
- [Developer Kit User Guide](#)
- [L4T Documentation](#)
- [Module Datasheet](#)
- [OEM Design Guide](#)



- [Jetson Nano Developer Kit](#)
- [Getting Started](#)
- [Developer Kit User Guide](#)
- [L4T Documentation](#)

JETSON NANO		JETSON TX2	JETSON XAVIER NX	JETSON AGX XAVIER
GPU	128 Core Maxwell 0.5 TFLOPs (FP16)	256 Core Pascal 1.3 TFLOPS (FP16)	384 Core Volta 21 TOPs (INT8)	512 Core Volta + NVDLA 10 TFLOPS (FP16) 32 TOPS (INT8)
CPU	4 core ARM A57	6 core Denver and A57 (2x) 2MB L2	6 core Carmel ARM CPU (3x) 2MB L2 + 4MB L3	8 core Carmel ARM CPU (4x) 2MB L2 + 4MB L3
Memory	4 GB 64-bit LPDDR4 25.6 GB/s	Up to 8 GB 128b LPDDR4 58 GB/s	8 GB 128-bit LPDDR4x 51.2 GB/s	Up to 32GB 256-bit LPDDR4x 137 GB/s
Storage	16 GB eMMC	Up to 32 GB eMMC	16 GB eMMC	32 GB eMMC
Encode	4K @ 30 (H.265)	4K @ 60 (H.265)	2x 4K @ 30 (H.265)	4x 4K @ 60 (H.265)
Decode	4K @ 60 (H.265)	2x 4K @ 60 (H.265)	2x 4K @ 60 (H.265)	6x 4K @ 60 (H.265)
Camera	12 (3x4 or 4x2) MIPI CSI-2 D-PHY 1.1 lanes (18 Gbps)	12 lanes MIPI CSI-2 D-PHY 1.2 (30 Gbps) C-PHY (41 Gbps)	12 lanes (3x4 or 6x2) MIPI CSI-2 D-PHY 1.2 (30 Gbps)	16 lanes MIPI CSI-2 8 lanes SLVS-EC D-PHY (40 Gbps) C-PHY (59 Gbps)
Mechanical	69.6mm x 45mm 260 pin edge connector	87mm x 50mm 400 pin connector	69.6mm x 45mm 260 pin edge connector	100mm x 87mm 699 pin connector
Software	JetPack SDK - Unified software release across all Jetson products			

JETSON SOFTWARE



AGENDA

Jetson product and marketing

Why we need cross compile

How to setup the cross-compile environment

1. Config directly via the environment variable.
 2. Config in Makefile.
 3. Build your simulation.
-

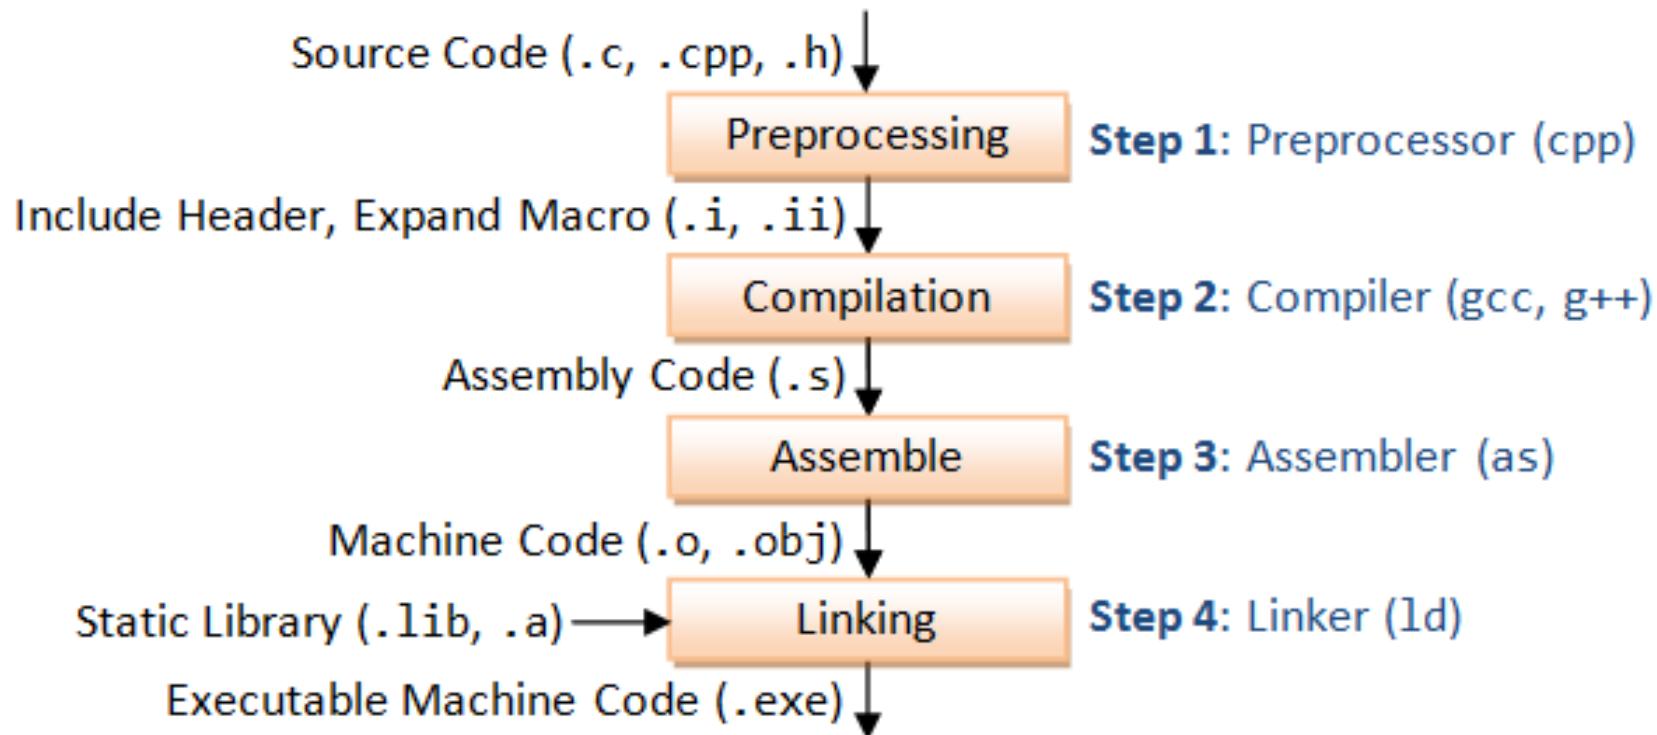
Compile on X86 host and execute on Jetson

1. Cross compile Linux kernel.
2. Cross compile CUDA project.
3. Cross compile Multi Media API based application.
4. Cross compile Deepstream.
5. Cross compile Caffe
6. Cross compile OpenCV.

Where can I get the prebuild version and how to build on jetson?

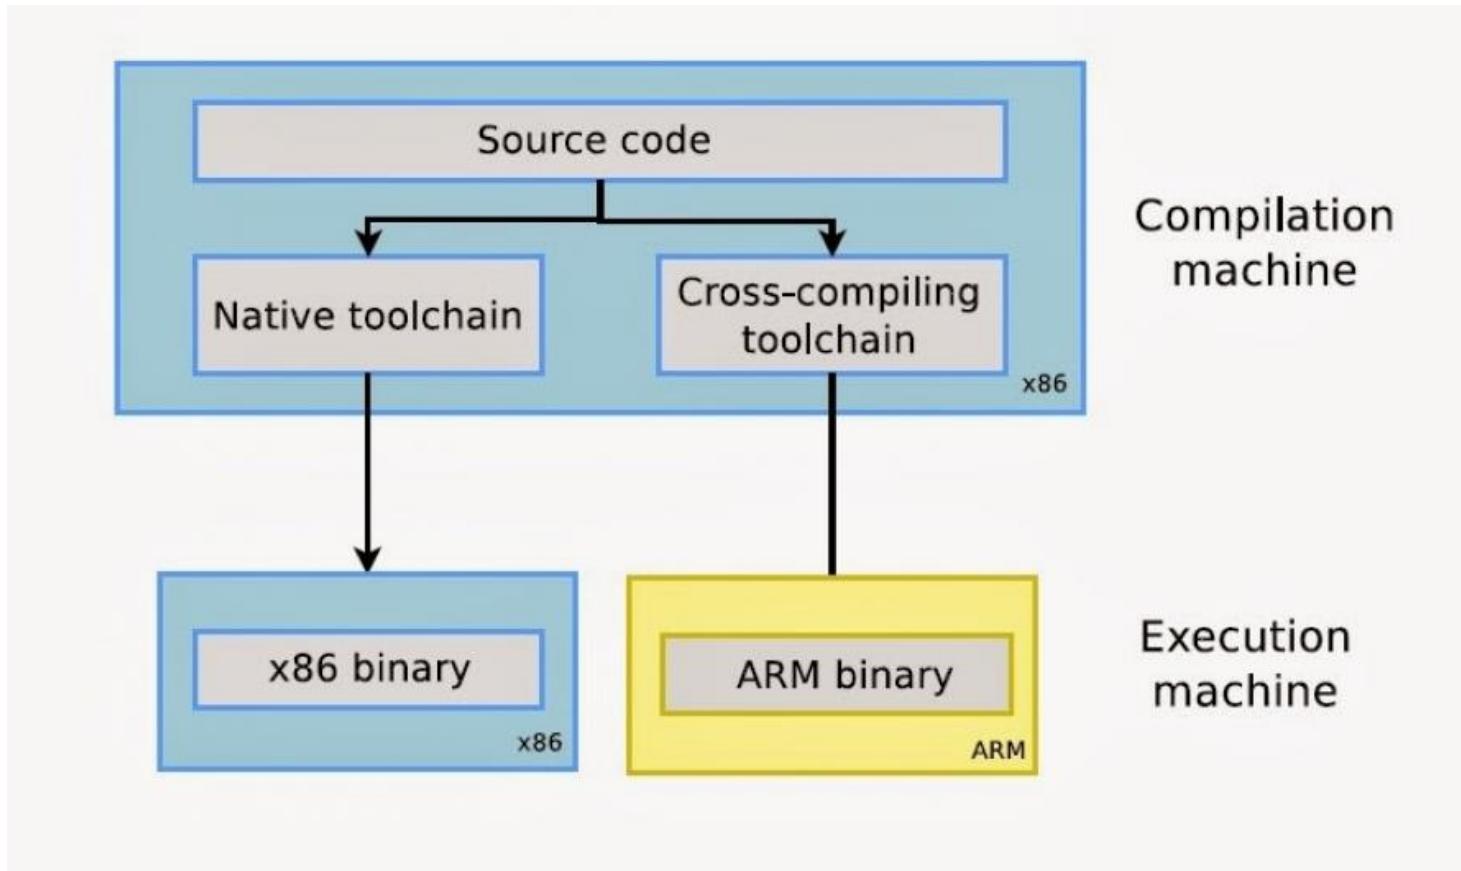
ARCHITECTURES & INSTRUCTION SETS

How compile works



ARCHITECTURES & INSTRUCTION SETS

Native compile and cross copile



ARCHITECTURES & INSTRUCTION SETS

What can we benefit from cross compile

- Speed
Target platform usually designed to have low cost and low power consumption. So the speed of the CPU will be not good at compiling.
- Capability
Target platform have less storage and less memory to sustain running a big project.
- Availability
Not all the system for target platform has the resources to support compile natively.
- Flexibility
A complete build environment need lots of the software to support.

AGENDA

Jetson product and marketing

Why we need cross compile

How to setup the cross-compile environment

1. Config directly via the environment variable.
 2. Config in Makefile.
 3. Build your simulation.
-

Compile on X86 host and execute on Jetson

1. Cross compile Linux kernel.
2. Cross compile CUDA project.
3. Cross compile Multi Media API based application.
4. Cross compile Deepstream.
5. Cross compile Caffe
6. Cross compile OpenCV.

Where can I get the prebuild version and how to build on jetson?

SETUP THE CROSS COMPILE

Dowload the toolchain

NVIDIA® specifies the Linaro® gcc 7.3.1 2018.05 aarch64 toolchain for:

- Cross-compiling applications to run on Jetson Linux Driver Package (L4T rel-32.
- Cross-compiling code in the L4T rel-32 source release.

This topic describes how to obtain this toolchain.

Toolchain Information

The toolchain contains the following components:

- GCC version: 7.3.1
- Binutils version: 2.28.2.20170706
- Glibc version: 2.25

Downloading the Toolchain

Download the pre-built toolchain binaries from:

http://releases.linaro.org/components/toolchain/binaries/7.3-2018.05/aarch64-linux-gnu/gcc-linaro-7.3.1-2018.05-x86_64_aarch64-linux-gnu.tar.xz

Extracting the Toolchain

Execute the following commands to extract the toolchain:

```
$ mkdir $HOME/l4t-gcc  
$ cd $HOME/l4t-gcc  
$ tar xf gcc-linaro-7.3.1-2018.05-x86_64_aarch64-linux-gnu.tar.xz
```

SETUP THE CROSS COMPILE

Options 1:

```
export CROSS_ROOT=/home/alan/rootfs-nano  
export CROSS_COMPILE = /home/alan/toolchain/gcc-4.8.5-aarch64/bin  
export CC = ${CROSS_COMPILE}/aarch64-unknown-linux-gnu-gcc  
export CXX = ${CROSS_COMPILE}/aarch64-unknown-linux-gnu-g++  
export LD = ${CROSS_COMPILE}/aarch64-unknown-linux-gnu-ld  
export AR = ${CROSS_COMPILE}/aarch64-unknown-linux-gnu-ar  
export AS = ${CROSS_COMPILE}/aarch64-unknown-linux-gnu-as  
export RANLIB = ${CROSS_COMPILE}/aarch64-unknown-linux-gnu-ranlib  
export NVCC = ${CROSS_ROOT}/usr/local/cuda/bin/nvcc  
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$SYSROOT/usr/lib/aarch64-linux-gnu:$SYSROOT/lib/aarch64-linux-gnu:$SYSROOT/lib  
export LC_ALL=C  
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:/home/alan/toolchain/gcc-4.8.5-aarch64/aarch64-unknown-linux-gnu/lib64:/home/alan/toolchain/gcc-4.8.5-aarch64/aarch64-unknown-linux-gnu/sysroot/lib
```

Options 2:

```
SYSROOT=/mnt
#export CROSS_COMPILE = /home/alan/toolchain/gcc-4.8.5-aarch64/bin
CROSS_ROOT=/home/green/Downloads/toolchain/aarch64
CROSS_COMPILE=${CROSS_ROOT}/bin
CC=${CROSS_COMPILE}/aarch64-unknown-linux-gnu-gcc
CXX=${CROSS_COMPILE}/aarch64-unknown-linux-gnu-g++
LD=${CROSS_COMPILE}/aarch64-unknown-linux-gnu-ld
AR=${CROSS_COMPILE}/aarch64-unknown-linux-gnu-ar
AS=${CROSS_COMPILE}/aarch64-unknown-linux-gnu-as
RANLIB=${CROSS_COMPILE}/aarch64-unknown-linux-gnu-ranlib
NVCC=${SYSROOT}/usr/local/cuda/bin/nvcc

CFLAGS:= -Wall -Werror -std=c++11 -fPIC -Wno-error=deprecated-declarations
CFLAGS+= -I${SYSROOT}/usr/include/aarch64-linux-gnu -I${SYSROOT}/usr/include -I${SYSROOT}/usr/local/cuda-10.0/targets/aarch64-linux/include \
I../../includes

LIBS:= -L/home/green/Downloads/toolchain/aarch64/aarch64-unknown-linux-gnu/sysroot/usr/lib /home/green/Downloads/toolchain/aarch64/aarch64- \
unknown-linux-gnu/sysroot/usr/lib/libc_nonshared.a -lvinfer -lnvparsers -L${SYSROOT}/usr/local/cuda/lib64 -L${SYSROOT}/usr/lib/aarch64-linux- \
gnu -lcudart -lcublas /home/green/Downloads/toolchain/aarch64/aarch64-unknown-linux-gnu/sysroot/usr/lib/libc_nonshared.a
LFLAGS:= -Wl,--start-group ${LIBS} -Wl,--end-group

SRCFILES:= nvdsparsebbox_ssd.cpp nvdsiplugin_ssd.cpp
TARGET_LIB:= libnvdsinfer_custom_impl_ssd.so

all: $(TARGET_LIB)

$(TARGET_LIB) : $(SRCFILES)
    $(CC) -o $@ $^ $(CFLAGS) $(LFLAGS)

clean:
    rm -rf $(TARGET_LIB)
```

Setup the NFS on Jetson

```
$ sudo apt update  
$ sudo apt install nfs-kernel-server  
$ sudo vim /etc/exports  
    / *(rw,sync,no_subtree_check, no_root_squash)  
$ sudo exportfs -a  
$ sudo systemctl restart nfs-kernel-server.service  
$ sudo exportfs  
    /      (world)
```

Setup the NFS on Host

```
$ sudo apt update
```

```
$ sudo apt install nfs-common
```

```
$ sudo mkdir -p /mnt/rootfs
```

```
$ sudo /etc/fstab
```

```
192.168.55.1:/ /mnt/rootfs nfs defaults 0 0
```

```
$ sudo mount -t nfs 192.168.55.1:/ /mnt/rootfs
```

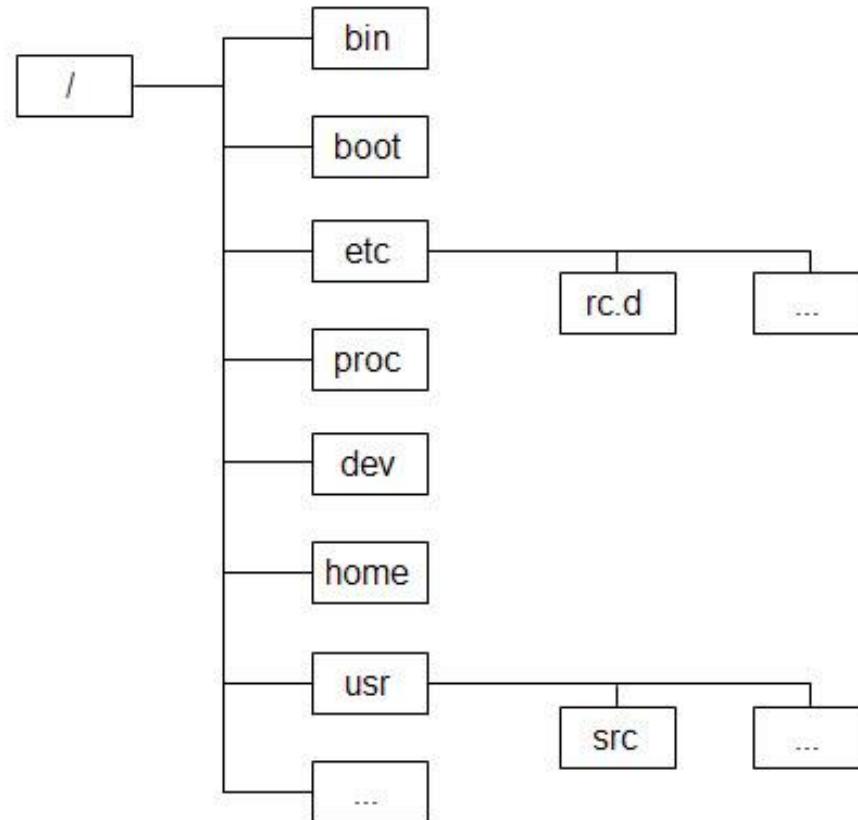
```
$ ls /mnt/rootfs/
```

```
green@C ~ $ ls /mnt/rootfs/
README.txt bin boot dev etc home lib lost+found media mnt opt proc root run sbin snap srv sys tmp usr var
```

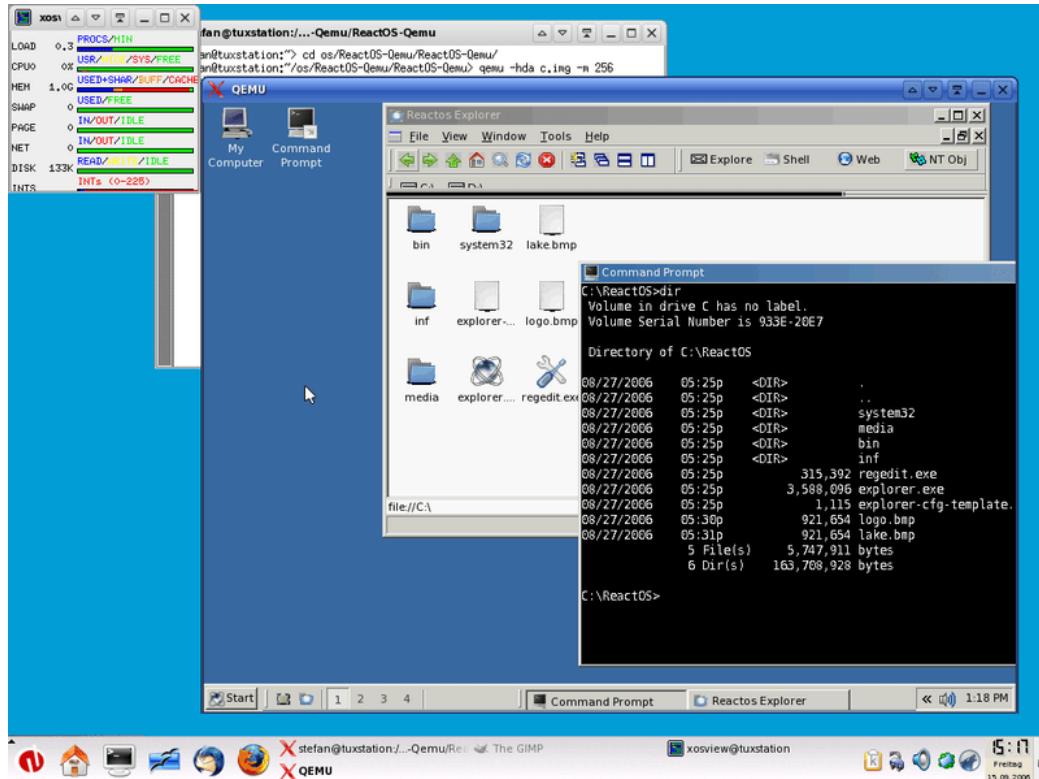
Chroot & Qeum

Options 3:

- 系统修复
- 系统启动时切换根目录，引导系统启动
- 特权分离



Chroot & Qemu



```
[test@donizetti ~]$ qemu-arm ./ls --color /
bin  etc  lib64  mnt  root  srv      system-upgrade-root  var
boot  home  lost+found  opt  run  sys
dev   lib   media  proc  sbin  system-upgrade  tmp
                proc
[test@donizetti ~]$ uname -a
Linux donizetti 4.6.7-300.fc24.x86_64 #1 SMP Wed Aug 17 18:48:43 UTC 2016 x86_64
x86_64 x86_64 GNU/Linux
[test@donizetti ~]$ file ./ls
./ls: ELF 32-bit LSB executable, ARM, EABI5 version 1 (SYSV), dynamically linked
, interpreter /lib/ld-linux-armhf.so.3, for GNU/Linux 3.0.0, stripped
[test@donizetti ~]$
```

<https://www.qemu.org/>

AGENDA

Jetson product and marketing

Why we need cross compile

How to setup the cross-compile environment

1. Config directly via the environment variable.
 2. Config in Makefile.
 3. Build your simulation.
-

Compile on X86 host and execute on Jetson

1. Cross compile Linux kernel.
2. Cross compile CUDA project.
3. Cross compile Multi Media API based application.
4. Cross compile Deepstream.
5. Cross compile Caffe
6. Cross compile OpenCV.

Where can I get the prebuild version and how to build on jetson?

Cross compile the Kernel

```
green@C ~/nvidia/nvidia_sdk $ ./source_sync.sh -t
Terminating... wrong switch: -t
Use: source_sync.sh [options]
Available general options are,
  -h      : help
  -e      : exit on sync error
  -d [DIR] : root of source is DIR
  -t [TAG] : Git tag that will be used to sync all the sources

By default, all sources are downloaded.
Only specified sources are downloaded, if one or more of the following options are mentioned.

  -k [TAG]: Download kernel/kernel-4.9 source and optionally sync to TAG
            and download kernel/nvgpu source and sync to the same TAG
            and download kernel/nvidia source and sync to the same TAG
            and download hardware/nvidia/soc/t18x source and sync to the same TAG
            and download hardware/nvidia/platform/tegra/common source and sync to the same TAG
            and download hardware/nvidia/platform/t18x/common source and sync to the same TAG
            and download hardware/nvidia/platform/t18x/quill source and sync to the same TAG
            and download hardware/nvidia/soc/t210 source and sync to the same TAG
            and download hardware/nvidia/platform/t210/common source and sync to the same TAG
            and download hardware/nvidia/platform/t210/jetson source and sync to the same TAG
            and download hardware/nvidia/platform/t210/porg source and sync to the same TAG
            and download hardware/nvidia/soc/t19x source and sync to the same TAG
            and download hardware/nvidia/platform/t19x/common source and sync to the same TAG
            and download hardware/nvidia/platform/t19x/galen/bct source and sync to the same TAG
            and download hardware/nvidia/platform/t19x/galen/kernel-dts source and sync to the same TAG
            and download hardware/nvidia/soc/tegra source and sync to the same TAG
  -u [TAG]: Download u-boot source and optionally sync to TAG
```

https://docs.nvidia.com/jetson/l4t-multimedia/cross_platform_support.html

Cross compile the Kernel

```
$ green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ git tag -l
```

```
green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ git tag -l
tegra-l4t-r31.0.1
tegra-l4t-r31.0.2
tegra-l4t-r31.1
tegra-l4t-r32.1
tegra-l4t-r32.1.update-01
tegra-l4t-r32.1.update-1
tegra-l4t-r32.2.0
tegra-l4t-r32.2.1
tegra-l4t-r32.2.2
tegra-l4t-r32.2.3-1
tegra-l4t-r32.3.1
```

<https://docs.nvidia.com/jetson/l4t/index.html>

Cross compile the Kernel

```
$ green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ ./source_sync.sh -t tegra-l4t-r32.3.1
```

```
green@C ~/nvidia/nvidia_sdk $ ./source_sync.sh -t tegra-l4t-r32.3.1
Directory for kernel/kernel-4.9, /home/green/nvidia/nvidia_sdk/sources/kernel/kernel-4.9, already exists!
Syncing up with tag tegra-l4t-r32.3.1...
Switched to a new branch 'mybranch_2020-04-16-1587032356'
/home/green/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 source sync'ed to tag tegra-l4t-r32.3.1 successfully!

Directory for kernel/nvgpu, /home/green/nvidia/nvidia_sdk/sources/kernel/nvgpu, already exists!
Syncing up with tag tegra-l4t-r32.3.1...
Switched to a new branch 'mybranch_2020-04-16-1587032357'
/home/green/nvidia/nvidia_sdk/sources/kernel/nvgpu source sync'ed to tag tegra-l4t-r32.3.1 successfully!

Directory for kernel/nvidia, /home/green/nvidia/nvidia_sdk/sources/kernel/nvidia, already exists!
Syncing up with tag tegra-l4t-r32.3.1...
Switched to a new branch 'mybranch_2020-04-16-1587032358'
/home/green/nvidia/nvidia_sdk/sources/kernel/nvidia source sync'ed to tag tegra-l4t-r32.3.1 successfully!

Directory for hardware/nvidia/soc/t18x, /home/green/nvidia/nvidia_sdk/sources/hardware/nvidia/soc/t18x, already exists!
Syncing up with tag tegra-l4t-r32.3.1...
Switched to a new branch 'mybranch_2020-04-16-1587032359'
/home/green/nvidia/nvidia_sdk/sources/hardware/nvidia/soc/t18x source sync'ed to tag tegra-l4t-r32.3.1 successfully!
```

<https://docs.nvidia.com/jetson/l4t/index.html>

Cross compile the Kernel

```
green@C ~/nvidia/nvidia_sdk/sources $ ls  
hardware kernel u-boot  
green@C ~/nvidia/nvidia_sdk/sources $ cd kernel/kernel-4.9/  
green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $  
  
export ARCH=arm64  
export CROSS_COMPILE=${HOME}/Downloads/toolchain/gcc-linaro-7.3.1-2018.05-x86_64_aarch64-linux-  
gnu/bin/aarch64-linux-gnu-  
  
export LOCALVERSION=-tegra
```

nvidia@nvidia-desktop:/usr/local/cuda/samples\$ uname -a
Linux nvidia-desktop 4.9.140-tegra #1 SMP PREEMPT Mon Dec 9 22:52:02 PST 2019 aarch64 aarch64 aarch64
GNU/Linux

\$ wget http://releases.linaro.org/components/toolchain/binaries/7.3-2018.05/aarch64-linux-gnu/gcc-linaro-7.3.1-2018.05-x86_64_aarch64-linux-gnu.tar.xz

Get the config file

Jetson:

```
nvidia@nvidia-desktop:/usr/local/cuda/samples$ sudo cp /proc/config.gz /
[sudo] password for nvidia:
nvidia@nvidia-desktop:/usr/local/cuda/samples$
```

X86 Host:

```
green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ sudo gzip -d /mnt/rootfs/config.gz
green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ cp /mnt/rootfs/config .
'./config' -> './config'
removed './config'
green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ cp config .config
'.config' -> '.config'
green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ make menuconfig
scripts/kconfig/mconf Kconfig
```

Make -j8

<https://docs.nvidia.com/jetson/l4t/index.html>

KERNEL CROSS COMPILE

```
green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ make Image
CHK  include/config/kernel.release
CHK  include/generated/uapi/linux/version.h
CHK  include/generated/utsrelease.h
CHK  include/generated/bounds.h
CHK  include/generated/timeconst.h
CHK  include/generated/asm-offsets.h
CALL  scripts/checksyscalls.sh
CHK  scripts/mod/devicetable-offsets.h
CHK  include/generated/compile.h
CHK  kernel/config_data.h
green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ ls arch/arm64/boot/
.Image.cmd    .Image.gz.cmd   .gitignore   .zImage.cmd   Image      Image.gz      Makefile      dts/      install.sh      zImage
green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ ls arch/arm64/boot/Image
arch/arm64/boot/Image
green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ make dtbs
CHK  scripts/mod/devicetable-offsets.h
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/jetson/kernel-dts/tegra210-jetson-cv-base-p2597-2180-a00.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/jetson/kernel-dts/tegra210-jetson-cv-p2597-2180-a00.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/jetson/kernel-dts/tegra210-jetson-cv-p2597-2180-imx274-hdmi.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/jetson/kernel-dts/tegra210-jetson-cv-p2597-2180-a00-ao0-1080p-edp.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/jetson/kernel-dts/tegra210-jetson-tx1-p2597-2180-a01-android-devkit.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/jetson/kernel-dts/tegra210-jetson-tx1-p2597-2180-a01-devkit.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/jetson/kernel-dts/tegra210-jetson-tx1-p2597-2180-a01-imx274-dp-hdmi.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/jetson/kernel-dts/tegra210-jetson-e-base-p2595-0000-a00.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/jetson/kernel-dts/tegra210-jetson-e-p2595-0000-a00-00.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/jetson/kernel-dts/tegra210-jetson-tx1-p2597-2180-a02-devkit-24x7.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/jetson/kernel-dts/tegra210-jetson-tx1-p2597-2180-a01-devkit_as_TM660M.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/porg/kernel-dts/tegra210-p3448-0000-p3449-0000-a00.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/porg/kernel-dts/tegra210-p3448-0000-p3449-0000-a01.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/porg/kernel-dts/tegra210-p3448-0000-p3449-0000-a02.dtb
DTC  arch/arm64/boot/dts/_ddot/_ddot/_ddot/_ddot/_hardware/nvidia/platform/t210/porg/kernel-dts/tegra210-p3448-0002-p3449-0000-a02.dtb
```

```
green@C ~/nvidia/nvidia_sdk/sources/kernel/kernel-4.9 $ cp arch/arm64/boot/Image /mnt/rootfs/boot/
```

Cross compile Multi Media API

```
export TARGET_ROOTFS=/mnt/rootfs

sudo ln -sf /mnt/rootfs/usr/lib/aarch64-linux-gnu /usr/lib/aarch64-linux-gnu

sudo ln -sf /usr/lib/aarch64-linux-gnu/libpthread.so /lib/libpthread.so.0

sudo ln -sf /mnt/rootfs/lib/aarch64-linux-gnu /lib/aarch64-linux-gnu

sudo ln -sf /mnt/rootfs/lib/aarch64-linux-gnu/libc.so.6 /lib/libc.so.6

sudo ln -sf /mnt/rootfs/usr/lib/aarch64-linux-gnu/libc_nonshared.a /usr/lib/libc_nonshared.a

sudo ln -sf /lib/aarch64-linux-gnu/ld-2.27.so /lib/ld-linux-aarch64.so.1
```

<https://developer.nvidia.com/nvidia-jetson-linux-multimediaapireference>

Cross compile Multi Media API

```
green@C ~/Downloads/nvidia/sdkm_downloads/tegra_multimedia_api $ git diff
diff --git a/samples/Rules.mk b/samples/Rules.mk
index 7977ab1..6dd362d 100644
--- a/samples/Rules.mk
+++ b/samples/Rules.mk
@@ -76,7 +76,7 @@ NM      = $(AT) $(CROSS_COMPILE)nm
 STRIP   = $(AT) $(CROSS_COMPILE)strip
 OBJCOPY = $(AT) $(CROSS_COMPILE)objcopy
 OJDUMP  = $(AT) $(CROSS_COMPILE)objdump
-NVCC    = $(AT) $(CUDA_PATH)/bin/nvcc -ccbin $(filter-out $(AT), $(CPP))
+NVCC    = $(AT) $(TARGET_ROOTFS)$(CUDA_PATH)/bin/nvcc -ccbin $(filter-out $(AT), $(CPP))

# Specify the logical root directory for headers and libraries.
ifeq ($(TARGET_ROOTFS),)
@@ -96,7 +96,8 @@ CPPFLAGS += -std=c++11 \
-I"$(ALGO_TRT_DIR)" \
-I"$(TARGET_ROOTFS)/$(CUDA_PATH)/include" \
-I"$(TARGET_ROOTFS)/usr/include/$(TEGRA_ARMABI)" \
- -I"$(TARGET_ROOTFS)/usr/include/libdrm"
+ -I"$(TARGET_ROOTFS)/usr/include/libdrm" \
+ -I"$(TARGET_ROOTFS)/usr/include/opencv4"

# All common dependent libraries
LDFLAGS += \
```

<https://developer.nvidia.com/nvidia-jetson-linux-multimediaapireference>

Cross compile Multi Media API

```
CUDA_PATH      := $(TOP_DIR)/samples/common/algorithm/trt
TOP_DIR
CLASS_DIR
ALGO_CUDA_DIR
ALGO_TRT_DIR
CROSS_COMPILE
CROSS_COMPILE
AS            := $(AT) $(CROSS_COMPILE)as
LD            := $(AT) $(CROSS_COMPILE)ld
CC            := $(AT) $(CROSS_COMPILE)gcc
CPP           := $(AT) $(CROSS_COMPILE)g++
AR            := $(AT) $(CROSS_COMPILE)ar
NM            := $(AT) $(CROSS_COMPILE)nm
STRIP          := $(AT) $(CROSS_COMPILE)strip
OBJCOPY        := $(AT) $(CROSS_COMPILE)objcopy
OBJDUMP        := $(AT) $(CROSS_COMPILE)objdump
NVCC          := $(AT) $(TARGET_ROOTFS)$(CUDA_PATH)/bin/nvcc -ccbin $(filter-out $(AT), $(CPP))

# Specify the logical root directory for headers and libraries.
ifeq ($(TARGET_ROOTFS),)
CPPFLAGS += --sysroot=$(TARGET_ROOTFS)
LDFLAGS += \
    -Wl,-rpath-link=$(TARGET_ROOTFS)/lib/$(TEGRA_ARMABI) \
    -Wl,-rpath-link=$(TARGET_ROOTFS)/usr/lib/$(TEGRA_ARMABI) \
    -Wl,-rpath-link=$(TARGET_ROOTFS)/usr/lib/$(TEGRA_ARMABI)/tegra \
    -Wl,-rpath-link=$(TARGET_ROOTFS)/$(CUDA_PATH)/lib64
endif

# All common header files
CPPFLAGS += -std=c++11 \
    -I$(TOP_DIR)/include \
    -I$(TOP_DIR)/include/libjpeg-8b \
    -I$(ALGO_CUDA_DIR) \
    -I$(ALGO_TRT_DIR) \
    -I$(TARGET_ROOTFS)$(CUDA_PATH)/include \
    -I$(TARGET_ROOTFS)/usr/include/$(TEGRA_ARMABI) \
    -I$(TARGET_ROOTFS)/usr/include/libdrm \
    -I$(TARGET_ROOTFS)/usr/include/opencv4

# All common dependent libraries
LDFLAGS += \
    -lpthread -lv4l2 -lEGL -lGLESv2 -lX11 \
    -lnvbuf_utils -lnvjpeg -lnvosd -ldr \
    -lcuda -lcudart \
    -lnvinfer -lnvparsers \
    -L$(TARGET_ROOTFS)$(CUDA_PATH)/lib64 \
    -L$(TARGET_ROOTFS)/usr/lib/$(TEGRA_ARMABI) \
    -L$(TARGET_ROOTFS)/usr/lib/$(TEGRA_ARMABI)/tegra
```

<https://developer.nvidia.com/nvidia-jetson-linux-multimediaapireference>

Cross compile Multi Media API

To run

```
$ ./video_dec_trt 2 ../../data/Video/sample_outdoor_car_1080p_10fps.h264  
../../data/Video/sample_outdoor_car_1080p_10fps.h264 H264 --trt-deployfile  
../../data/Model/resnet10/resnet10.prototxt --trt-modelfile ../../data/Model/resnet10/resnet10.caffemodel --trt-  
mode 0
```

```
nvidia@nvidia-desktop:~/tegra_multimedia_api/samples/backend$ ./backend 1 ../../data/Video/sample_outdoor_car_1080p_10fps.h264 H264 \ --trt-deleNet_one_class/GoogleNet_modified_oneClass_halfHD.caffemodel \
nvidia@nvidia-desktop:~/tegra_multimedia_api/samples/backend$ ./backend 1 ../../data/Video/sample_outdoor_car_1080p_10fps.h264 H264 \
>     --trt-deployfile ../../data/Model/GoogleNet_one_class/GoogleNet_modified_oneClass_halfHD.prototxt \
>     --trt-modelfile ../../data/Model/GoogleNet_one_class/GoogleNet_modified_oneClass_halfHD.caffemodel \
>     --trt-mode 0 --trt-proc-interval 1 -fps 10
Net has batch_size, channel, net_height, net_width: 1 3 540 960
mode has been set to 0 (using fp16)
outputs coverage
outputs bboxes
Applying generic optimizations to the graph for inference.
Original: 142 layers
After dead-layer removal: 142 layers
After scale fusion: 142 layers
Fusing conv1/7x7_s2 with conv1/relu_7x7
Fusing conv2/3x3_reduce with conv2/relu_3x3_reduce
Fusing conv2/3x3 with conv2/relu_3x3
Fusing inception_3a/1x1 with inception_3a/relu_1x1
Fusing inception_3a/3x3_reduce with inception_3a/relu_3x3_reduce
Fusing inception_3a/3x3 with inception_3a/relu_3x3
Fusing inception_3a/5x5_reduce with inception_3a/relu_5x5_reduce
Fusing inception_3a/5x5 with inception_3a/relu_5x5
Fusing inception_3a/pool_proj with inception_3a/relu_pool_proj
Fusing inception_3b/1x1 with inception_3b/relu_1x1
Fusing inception_3b/3x3_reduce with inception_3b/relu_3x3_reduce
```

<https://developer.nvidia.com/nvidia-jetson-linux-multimediaapireference>

Cross compile Multi Media API

Example

```
$ ./backend 1 ../../data/Video/sample_outdoor_car_1080p_10fps.h264 H264 --trt-deployfile  
../../data/Model/GoogleNet_one_class/GoogleNet_modified_oneClass_halfHD.prototxt --trt-modelfile  
../../data/Model/GoogleNet_one_class/GoogleNet_modified_oneClass_halfHD.caffemodel --trt-mode 0 --trt-proc-  
interval 1 -fps 10
```

```
nvidia@nvidia-desktop:~/tegra_multimedia_api/samples/04_video_dec_trt$ ./video_dec_trt 2 ../../data/Video/sample_outdoor_car_1080p_10fps.h264 ../../data/Video/  
--trt-modelfile ../../data/Model/resnet10/resnet10.caffemodel --trt-mode 0  
set deployfile: ../../data/Model/resnet10/resnet10.prototxt  
set modefile: ../../data/Model/resnet10/resnet10.caffemodel  
Net has batch_size, channel, net_height, net_width:2 3 368 640  
Using cached TRT model  
Deserialize required 2206394 microseconds.  
outputDim c 4 w 40 h 23  
outputDimsBBOX.c() 16 w 40 h 23  
Opening in BLOCKING MODE  
NvMMLiteOpen : Block : BlockType = 261  
Opening in BLOCKING MODE  
NvMMLiteOpen : Block : BlockType = 261  
NVMEDIA: Reading vendor.tegra.display-size : status: 6  
NVMEDIA: Reading vendor.tegra.display-size : status: 6  
NvMMLiteBlockCreate : Block : BlockType = 261  
Starting decoder capture loop thread  
NvMMLiteBlockCreate : Block : BlockType = 261  
Starting decoder capture loop thread  
Video Resolution: 1920x1080  
Video Resolution: 1920x1080  
Resolution change successful  
Input file read complete  
Resolution change successful
```

<https://developer.nvidia.com/nvidia-jetson-linux-multimediaapireference>

Cross compile Deepstream

```
export TARGET_ROOTFS=/mnt/rootfs
```

```
export CUDA_VER=10.0
```

```
" Press <F1> to display help or type <Esc> to quit
Makefile (/mnt/rootfs)
macro
  CUDA_VER
  AT
  AT
  TEGRA_ARMABI
  TARGET_ROOTFS
  CROSS_COMPILE
  CROSS_COMPILE
  AS
  LD
  CC
  CPP
  AR
  NM
  STRIP
  OBJCOPY
  OBJDUMP
  NVCC
  LFLAGS
  SRCFILES
  TARGET_LIB
  " Press <F1> to display help or type <Esc> to quit
  ifndef $(shell uname -m), aarch64
  endif
  CROSS_COMPILE :=
  else
  CROSS_COMPILE ?= aarch64-unknown-linux-gnu-
  endif
  AS      = $(AT) $(CROSS_COMPILE)as
  LD      = $(AT) $(CROSS_COMPILE)ld
  CC      = $(AT) $(CROSS_COMPILE)gcc
  CPP     = $(AT) $(CROSS_COMPILE)g++
  AR      = $(AT) $(CROSS_COMPILE)ar
  NM      = $(AT) $(CROSS_COMPILE)nm
  STRIP   = $(AT) $(CROSS_COMPILE)strip
  OBJCOPY = $(AT) $(CROSS_COMPILE)objcopy
  OBJDUMP = $(AT) $(CROSS_COMPILE)objdump
  NVCC    = $(AT) $(TARGET_ROOTFS)/usr/local/cuda/bin/nvcc -ccbin $(filter-out $(AT), $(CPP))
# Specify the logical root directory for headers and libraries.
ifeq ($(TARGET_ROOTFS),)
CPPFLAGS += --sysroot=$(TARGET_ROOTFS)
LDFLAGS += \
  -WL,-rpath-link=$(TARGET_ROOTFS)/lib/$(TEGRA_ARMABI) \
  -WL,-rpath-link=$(TARGET_ROOTFS)/usr/lib/$(TEGRA_ARMABI) \
  -WL,-rpath-link=$(TARGET_ROOTFS)/usr/lib/$(TEGRA_ARMABI)/tegra \
  -WL,-rpath-link=$(TARGET_ROOTFS)/usr/local/cuda/lib64
endif
#
#All common header files
CPPFLAGS += -std=gnu++11 -Wall -Werror -shared -fPIC -Wno-error=deprecated-declarations \
-I../../includes \
-I$(TARGET_ROOTFS)/usr/local/cuda/include" \
-I"$(TARGET_ROOTFS)/usr/include/$(TEGRA_ARMABI)" \
-I"$(TARGET_ROOTFS)/usr/include/libdrm" \
-I"$(TARGET_ROOTFS)/usr/includeopencv4"
#
LDFLAGS += \
-lcudart -lcublas \
-lnvinfer -lnvparsers \
-L"$(TARGET_ROOTFS)/usr/local/cuda/lib64" \
-L"$(TARGET_ROOTFS)/usr/lib/$(TEGRA_ARMABI)" \
-L"$(TARGET_ROOTFS)/usr/lib/$(TEGRA_ARMABI)/tegra"
#
LFLAGS:= -WL,--start-group $(LDFLAGS) -WL,--end-group
SRCFILES:= nvdsparsebbox_ssd.cpp nvdsiplugin_ssd.cpp
TARGET_LIB:= libnvdsinfer_custom_impl_ssd.so
#
$(TARGET_LIB) : $(SRCFILES)
    $(CC) -o $@ $^ $(CPPFLAGS) $(LFLAGS)
clean:
    rm -rf $(TARGET_LIB)
```

Pre-requisites for SSD

- Copy the model's label file "ssd_coco_labels.txt" from the data/ssd directory in TensorRT samples to this directory.

- Steps to generate the UFF model from ssd_inception_v2_coco TensorFlow frozen graph. These steps have been referred from TensorRT sampleUffSSD README:

1. Make sure TensorRT's uff-converter-tf package is installed.

2. Install tensorflow-gpu package for python:

For dGPU:

```
$ pip install tensorflow-gpu
```

For Jetson, refer to https://elinux.org/Jetson_Zoo#TensorFlow

3. Download and untar the ssd_inception_v2_coco TensorFlow trained model from http://download.tensorflow.org/models/object_detection/ssd_inception_v2_coco_2017_11_17.tar.gz

4. Navigate to the extracted directory and convert the frozen graph to uff:

```
$ cd ssd_inception_v2_coco_2017_11_17  
$ python /usr/lib/python2.7/dist-packages/uff/bin/convert_to_uff.py \  
frozen_inference_graph.pb -O NMS \  
-p /usr/src/tensorrt/samples/sampleUffSSD/config.py \  
-o sample_ssd_relu6.uff
```

5. Copy sample_ssd_relu6.uff to this directory.

Cross compile Deepstream

On the Host:

```
#cd /mnt/rootfs/home/nvidia/deepstream/deepstream-  
4.0/sources/objectDetector_SSD/nvdsinfer_custom_impl_ssd/ssd_inception_v2_coco_2017_11_17  
  
# python /usr/lib/python3.6/dist-packages/uff/bin/convert_to_uff.py  
/mnt/rootfs/home/nvidia/deepstream/deepstream-  
4.0/sources/objectDetector_SSD/nvdsinfer_custom_impl_ssd/ssd_inception_v2_coco_2017_11_17/frozen_infere  
nce_graph.pb -O NMS -p /usr/src/tensorrt/samples/sampleUffSSD/config.py -o  
/mnt/rootfs/home/nvidia/deepstream/deepstream-4.0/sources/objectDetector_SSD/sample_ssdlrelu6.uff  
  
# cd /mnt/rootfs/home/nvidia/deepstream/deepstream-4.0/sources/objectDetector_SSD/  
  
# cp /usr/src/tensorrt/data/ssd/ssd_coco_labels.txt .
```

On the Jetson:

```
# export DISPLAY=:0  
  
# deepstream-app -c deepstream_app_config_ssdl.txt
```

Cross compile DeepStream



Cross compile DeepStream

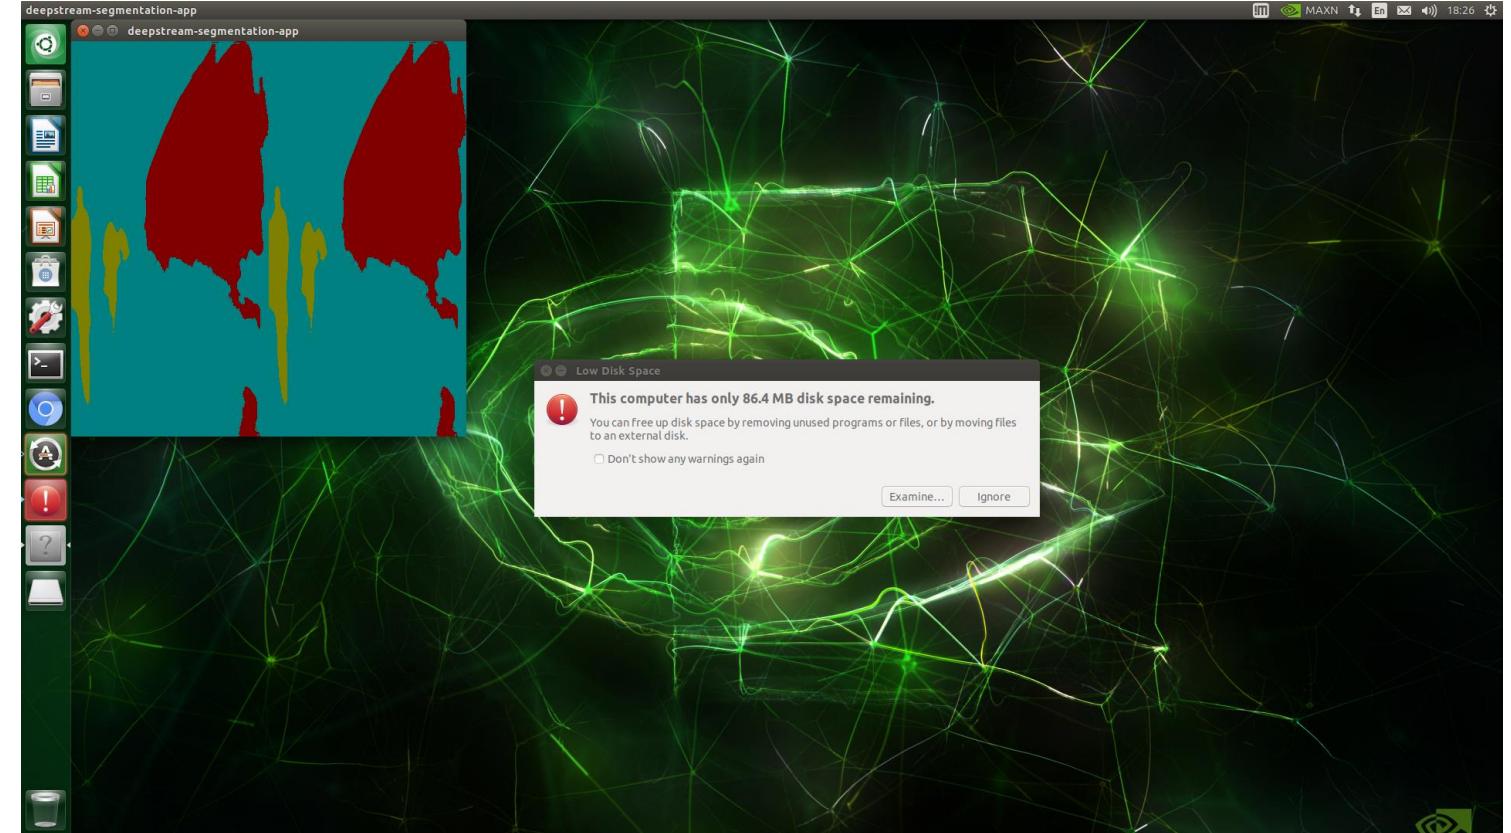
```
# cd  
/mnt/rootfs/home/nvidia/deepstream/deepstre  
am-4.0/sources/apps/sample_apps/deepstream-  
segmentation-test# make
```

```
Makefile (/mnt/rootfs/  
macro AT  
AT  
TEGRA_ARMABI  
TARGET_ROOTFS  
CROSS_COMPILE := aarch64-unknown-linux-gnu  
endif  
AS = $(AT) $(CROSS_COMPILE)as  
LD = $(AT) $(CROSS_COMPILE)ld  
CC = $(AT) $(CROSS_COMPILE)gcc  
CPP = $(AT) $(CROSS_COMPILE)g++  
AR = $(AT) $(CROSS_COMPILE)ar  
NM = $(AT) $(CROSS_COMPILE)nm  
STRIP = $(AT) $(CROSS_COMPILE)strip  
OBJCOPY = $(AT) $(CROSS_COMPILE)objcopy  
OBJDUMP = $(AT) $(CROSS_COMPILE)objdump  
NVCC = $(AT) $(TARGET_ROOTFS)/usr/local/cuda/bin/nvcc -ccbin $(filter-out $(AT), $(CPP))  
  
# Specify the logical root directory for headers and libraries.  
ifeq ($(TARGET_ROOTFS),)  
CFLAGS += --sysroot=$(TARGET_ROOTFS)  
LDFLAGS += \  
-Wl,-rpath-link=$(TARGET_ROOTFS)/lib/$(TEGRA_ARMABI) \  
-Wl,-rpath-link=$(TARGET_ROOTFS)/usr/lib/$(TEGRA_ARMABI) \  
-Wl,-rpath-link=$(TARGET_ROOTFS)/usr/lib/$(TEGRA_ARMABI)/tegra \  
-Wl,-rpath-link=$(TARGET_ROOTFS)/usr/local/cuda/lib64  
endif  
  
LIBS := deepstream-segmentation-app  
  
TARGET_DEVICE = $(shell gcc -dumpmachine | cut -f1 -d -)  
  
NVDS_VERSION:=4.0  
  
LIB_INSTALL_DIR?=/mnt/rootfs/opt/nvidia/deepstream/deepstream-$(NVDS_VERSION)/lib/ -L/mnt/rootfs/usr/lib/aarch64-linux-gnu $(LDFLAGS)  
  
ifeq ($(TARGET_DEVICE),aarch64)  
CFLAGS+= -DPLATFORM_TEGRA  
endif  
  
SRCS:= $(wildcard *.c)  
INCS:= $(wildcard *.h)  
PKGS:= gstreamer-1.0  
OBJS:= $(SRCS:.c=.o)  
CFLAGS+= -I../../..../includes  
CFLAGS+= -pthread -I$(TARGET_ROOTFS)/usr/include/gstreamer-1.0 -I$(TARGET_ROOTFS)/usr/include/glib-2.0 -I$(TARGET_ROOTFS)/usr/lib/aarch64-linux-gnu/glib-2.0/include  
LIBS:= -lgstreamer-1.0 -lgobject-2.0 -lplib-2.0  
LIBS+= -lm -L$(LIB_INSTALL_DIR) -lnvdsgst_helper -lnvdsgst_meta \  
-Wl,-rpath,$(LIB_INSTALL_DIR)  
all: $(APP)
```

Cross compile DeepStream

On Jetson

```
# ./deepstream-segmentation-app  
dstest_segmentation_config_semantic.txt  
sample_720p.mjpeg sample_720p.mjpeg
```



SETUP THE CROSS COMPILE

Option 1:

```
export TARGET_ROOTFS=/mnt/rootfs  
  
export CROSS_COMPILE=${HOME}/Downloads/toolchain/gcc-linaro-7.3.1-2018.05-x86_64_aarch64-linux-gnu/bin/aarch64-linux-gnu-  
  
export CC = ${CROSS_COMPILE}/gcc  
export CXX = ${CROSS_COMPILE}/g++  
export LD = ${CROSS_COMPILE}/ld  
export AR = ${CROSS_COMPILE}/ar  
export AS = ${CROSS_COMPILE}/as  
export RANLIB = ${CROSS_COMPILE}/ranlib  
export NVCC = ${TARGET_ROOTFS}/usr/local/cuda/bin/nvcc  
  
export LC_ALL=C
```

SETUP THE CROSS COMPILE

Option 2:

```
for i in /dev /dev/pts /proc /sys /run; do sudo mount -B $i /mnt$i; done
```

```
sudo mount -t proc /proc proc  
sudo mount --rbind /sys sys  
sudo mount --rbind /dev dev  
sudo mount --rbind /run run
```

```
sudo mount -o bind /dev dev  
sudo mount -o bind /sys sys  
sudo mount -t proc /proc proc  
sudo mount -o bind /run run
```

```
Sudo rm /etc/resolv.conf echo 'nameserver 8.8.4.4' | sudo tee -a /etc/resolv.conf
```

SETUP THE QEMU

```
for i in /dev /dev/pts /proc /sys /run; do sudo mount -B $i /mnt$i; done
```

```
sudo mount -o bind /dev dev
```

```
sudo mount -o bind /sys sys
```

```
sudo mount -t proc /proc proc
```

```
sudo mount -o bind /run run
```

```
sudo chroot .
```

<https://github.com/zhj-buffer/Cross-Compile-Jetson>

CUDA Cross Compile & Demo

```
$ cd /mnt/rootfs/usr/local/cuda/bin/  
  
$ ./cuda-install-samples-10.0.sh ~/  
Copying samples to /home/green/NVIDIA_CUDA-10.0_Samples now...  
Finished copying samples.  
green@C /mnt/rootfs/usr/local/cuda/bin $ █  
  
$ cd ~/NVIDIA_CUDA-10.0_Samples  
  
$ make -j8  
  
$ cp bin/aarch64/linux/release/ /mnt/rootfs/home/nvidia/ -rf
```

Verify on Jetson

```
nvidia@nvidia-desktop:/usr/local/cuda/samples/1_Utils...$ ./release/deviceQuery
/home/nvidia/release/deviceQuery Starting...

    CUDA Device Query (Runtime API) version (CUDART static linking)

Detected 1 CUDA Capable device(s)

Device 0: "Xavier"
    CUDA Driver Version / Runtime Version      10.0 / 10.0
    CUDA Capability Major/Minor version number: 7.2
    Total amount of global memory:              15815 MBytes (16583041024 bytes)
    ( 8) Multiprocessors, ( 64) CUDA Cores/MP:
    GPU Max Clock rate:                      1377 MHz (1.38 GHz)
    Memory Clock rate:                       1377 Mhz
    Memory Bus Width:                         256-bit
    L2 Cache Size:                           524288 bytes
    Maximum Texture Dimension Size (x,y,z):   1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
    Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
    Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
    Total amount of constant memory:           65536 bytes
    Total amount of shared memory per block:    49152 bytes
    Total number of registers available per block: 65536
    Warp size:                                32
    Maximum number of threads per multiprocessor: 2048
    Maximum number of threads per block:        1024
    Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
    Max dimension size of a grid size (x,y,z):  (2147483647, 65535, 65535)
    Maximum memory pitch:                     2147483647 bytes
    Texture alignment:                        512 bytes
    Concurrent copy and kernel execution:     Yes with 1 copy engine(s)
    Run time limit on kernels:                No
    Integrated GPU sharing Host Memory:       Yes
    Support host page-locked memory mapping: Yes
    Alignment requirement for Surfaces:       Yes
    Device has ECC support:                  Disabled
    Device supports Unified Addressing (UVA): Yes
    Device supports Compute Preemption:       Yes
    Supports Cooperative Kernel Launch:       Yes
    Supports MultiDevice Co-op Kernel Launch: Yes
    Device PCI Domain ID / Bus ID / location ID: 0 / 0 / 0
    Compute Mode:
        < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 10.0, CUDA Runtime Version = 10.0, NumDevs = 1
Result = PASS
nvidia@nvidia-desktop:/usr/local/cuda/samples/1_Utils...$
```

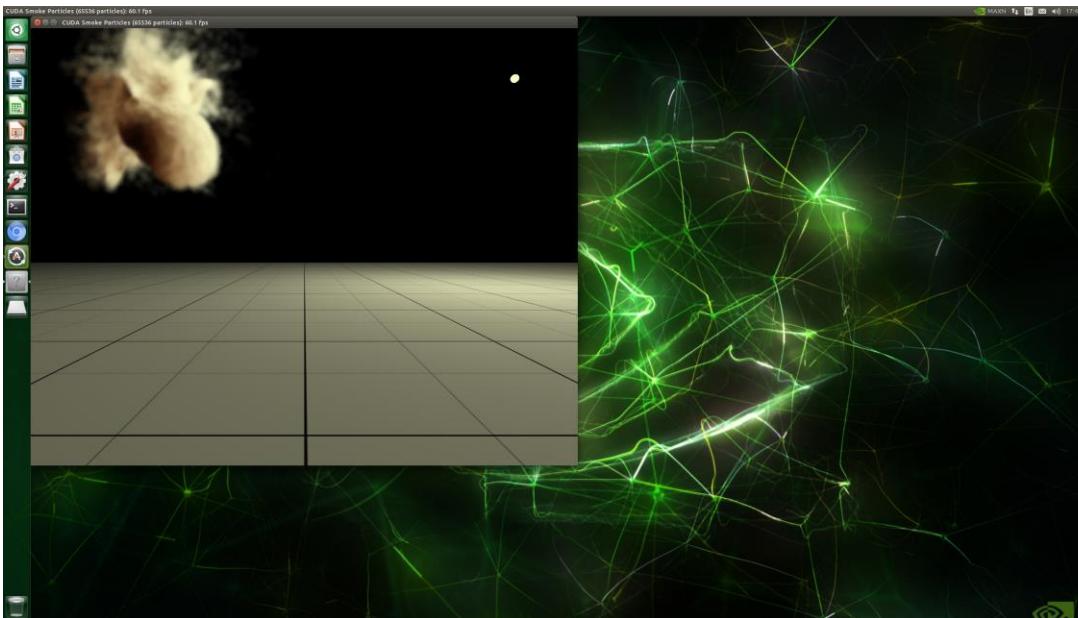
DeviceQuery

```
$ cd  
/usr/local/cuda/samples/1_Utilities/deviceQuer  
y  
$ sudo make  
$ ./deviceQuery
```

CUDA Capability 7.2

```
nvidia@nvidia-desktop:/usr/local/cuda/samples/1_Utilities/deviceQuery$ ./deviceQuery  
./deviceQuery Starting...  
  
CUDA Device Query (Runtime API) version (CUDART static linking)  
  
Detected 1 CUDA Capable device(s)  
  
Device 0: "Xavier"  
    CUDA Driver Version / Runtime Version      10.0 / 10.0  
    CUDA Capability Major/Minor version number: 7.2  
    Total amount of global memory:             15815 MBytes (16583041024 bytes)  
    ( 8) Multiprocessors, ( 64) CUDA Cores/MP:  
    GPU Max Clock rate:                      1377 MHz (1.38 GHz)  
    Memory Clock rate:                       1377 Mhz  
    Memory Bus Width:                        256-bit  
    L2 Cache Size:                           524288 bytes  
    Maximum Texture Dimension Size (x,y,z): 1D=(131072), 2D=(131072, 65536), 3D=(16384, 16  
384, 16384)  
    Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers  
    Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers  
    Total amount of constant memory:          65536 bytes  
    Total amount of shared memory per block:   49152 bytes  
    Total number of registers available per block: 65536  
    Warp size:                                32  
    Maximum number of threads per multiprocessor: 2048  
    Maximum number of threads per block:        1024  
    Max dimension size of a thread block (x,y,z): (1024, 1024, 64)  
    Max dimension size of a grid size (x,y,z): (2147483647, 65535, 65535)  
    Maximum memory pitch:                     2147483647 bytes  
    Texture alignment:                         512 bytes  
    Concurrent copy and kernel execution:     Yes with 1 copy engine(s)  
    Run time limit on kernels:                 No  
    Integrated GPU sharing Host Memory:       Yes  
    Support host page-locked memory mapping: Yes  
    Alignment requirement for Surfaces:       Yes  
    Device has ECC support:                   Disabled  
    Device supports Unified Addressing (UVA): Yes  
    Device supports Compute Preemption:       Yes  
    Supports Cooperative Kernel Launch:       Yes  
    Supports MultiDevice Co-op Kernel Launch: Yes  
    Device PCI Domain ID / Bus ID / location ID: 0 / 0 / 0  
    Compute Mode:  
        < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >  
  
deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 10.0, CUDA Runtime Version = 10.0, Num  
Devs = 1  
Result = PASS  
nvidia@nvidia-desktop:/usr/local/cuda/samples/1_Utilities/deviceQuery$
```

CUDA Cross Compile & Demo



```
nvidia@nvidia-desktop:/usr/local/cuda/samples$ ~/release/smokeParticles
CUDA Smoke Particles Starting...
NOTE: The CUDA Samples are not meant for performance measurements. Results may vary when GPU Boost is enabled.

Loaded './5_Simulations/smokeParticles/data/floortile.ppm', 256 x 256 pixels
GPU Device 0: "Xavier" with compute capability 7.2

^C
nvidia@nvidia-desktop:/usr/local/cuda/samples$
```

```
$ nvidia@nvidia-desktop:/usr/local/cuda/samples$ ~/release/smokeParticles
```

Compile Caffe & Test

Pre-request for Caffe

```
wget https://github.com/Kitware/CMake/releases/download/v3.17.1/cmake-3.17.1.tar.gz
```

```
apt-get install libssl-dev
```

```
./configure
```

```
Make -j16
```

```
Make install
```

```
/usr/local/bin/cmake
```

<https://github.com/zhj-buffer/Cross-Compile-Jetson>

<http://caffe.berkeleyvision.org/installation.html>

Compile Caffe & Test

Pre-request for Caffe

```
git clone https://github.com/BVLC/caffe.git
```

```
https://cmake.org/download/
```

```
sudo apt install libatlas-base-dev
```

```
sudo apt install libatlas-dev
```

```
sudo apt install libopenblas-dev
```

```
sudo apt install libsnavy-dev
```

```
sudo apt install libleveldb-dev
```

```
sudo apt install liblmdb-dev
```

```
apt install libboost-all-dev
```

```
apt install libgflags-dev
```

```
apt install libgoogle-glog-dev
```

```
sudo rm /etc/resolv.conf echo 'nameserver 8.8.4.4' | sudo tee -a /etc/resolv.conf
```

```
https://github.com/zj-buffer/Cross-Compile-Jetson
```

```
http://caffe.berkeleyvision.org/installation.html
```

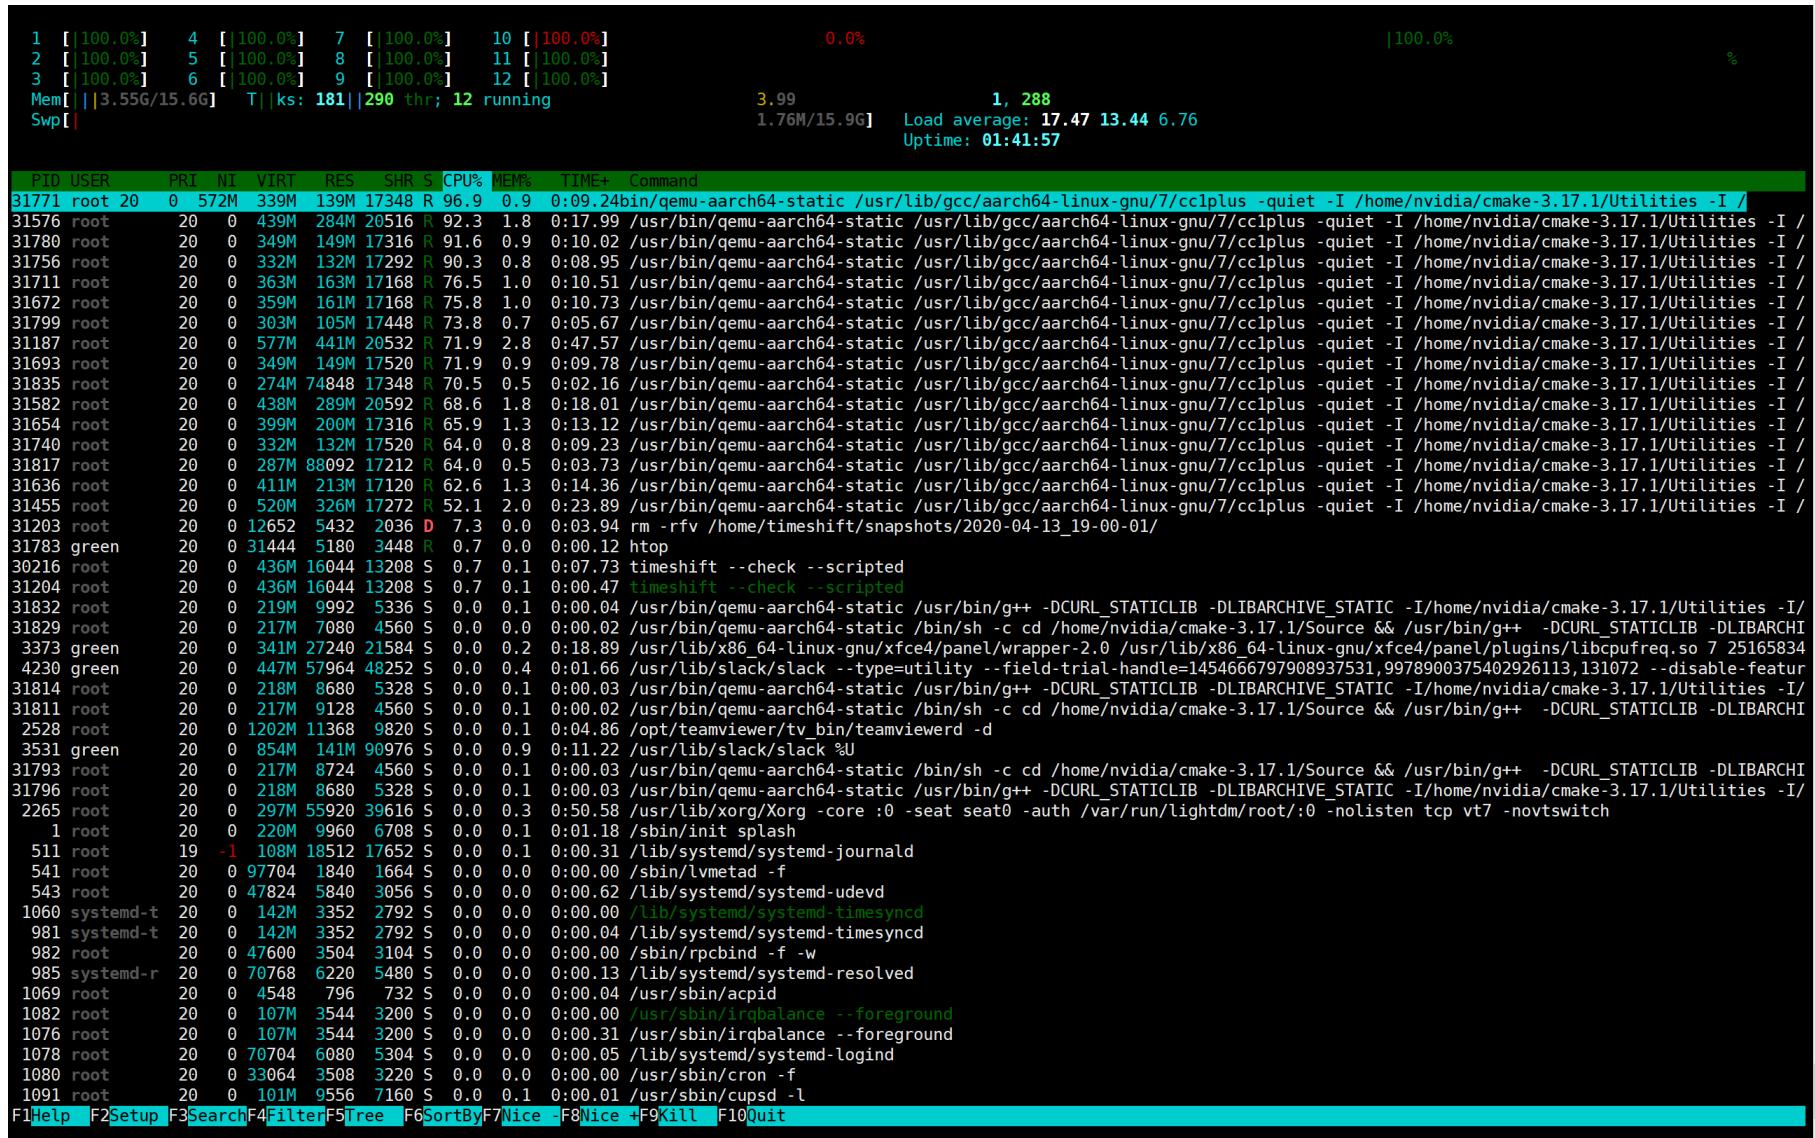
Compile Caffe & Test

Host:

```
cd caffe  
vim ../cmake/Cuda.cmake  
set(Caffe_known_gpu_archs "72")  
  
change CV_LOAD_IMAGE_COLOR to cv::IMREAD_COLOR  
change CV_LOAD_IMAGE_GRAYSCALE to cv::IMREAD_GRAYSCALE  
  
mkdir build  
  
/usr/local/bin/cmake ../  
  
make all  
  
make test  
  
make pycaffe  
  
cd /media/green/M0/rootfs-xavier/home/nvidia/caffe/build  
sudo cp lib /mnt/rootfs/home/nvidia/caffe/build/ -rf  
sudo cp ..//python/ /mnt/rootfs/home/nvidia/caffe/ -rf
```

Jetson:

```
export PYTHONPATH=/home/nvidia/caffe/python:$PYTHONPATH  
  
Python  
  
Import caffe
```



Compile Caffe & test

```
-- **** Caffe Configuration Summary ****
-- General:
--   Version      : 1.0.0
--   Git          : 1.0-136-g9b891540
--   System        : Linux
--   C++ compiler  : /usr/bin/c++
--   Release CXX flags : -O3 -DNDEBUG -fPIC -Wall -Wno-sign-compare -Wno-uninitialized
--   Debug CXX flags : -g -fPIC -Wall -Wno-sign-compare -Wno-uninitialized
--   Build type    : Release
--
--   BUILD_SHARED_LIBS : ON
--   BUILD_python     : ON
--   BUILD_matlab    : OFF
--   BUILD_docs       : ON
--   CPU_ONLY         : OFF
--   USE_OPENCV       : ON
--   USE_LEVELDB      : ON
--   USE_LMDB         : ON
--   USE_NCCL         : OFF
--   ALLOW_LMDB_NOLOCK : OFF
--   USE_HDF5         : ON
--
-- Dependencies:
--   BLAS          : Yes (Atlas)
--   Boost         : Yes (ver. 1.65)
--   glog          : Yes
--   gflags         : Yes
--   protobuf       : Yes (ver. 3.0.0)
--   lmdb          : Yes (ver. 0.9.21)
--   LevelDB        : Yes (ver. 1.20)
--   Snappy         : Yes (ver. ...)
--   OpenCV         : Yes (ver. 4.1.1)
--   CUDA          : Yes (ver. 10.0)
--
-- NVIDIA CUDA:
--   Target GPU(s)  : Auto
--   GPU arch(s)    : sm_20 sm_21 sm_30 sm_35 sm_50 sm_60 sm_61
--   cuDNN          : Yes (ver. 7.6.3)
--
-- Python:
--   Interpreter    : /usr/bin/python2.7 (ver. 2.7.17)
--   Libraries      : /usr/lib/aarch64-linux-gnu/libpython2.7.so (ver 2.7.17)
--   NumPy          : /usr/lib/python2.7/dist-packages/numpy/core/include (ver 1.13.3)
--
-- Documentation:
--   Doxygen        : No
--   config_file    :
--
-- Install:
--   Install path   : /home/nvidia/caffe/caffe/build/install
--
-- Configuring done
-- Generating done
-- Build files have been written to: /home/nvidia/caffe/caffe/build
root@C:/home/nvidia/caffe/caffe/build# make -j16
[ 0%] Running C++/Python protocol buffer compiler on /home/nvidia/caffe/caffe/src/caffe/proto/caffe.proto
Scanning dependencies of target caffeprotobuf
[ 1%] Building CXX object src/caffe/CMakeFiles/caffeprotobuf.dir/_/_/include/caffe/proto/caffe.pb.cc.o
[]
```

Compile Caffe & Test

```
-- Version      : 1.0.0
-- Git          : 1.0-136-g9b891540-dirty
-- System       : Linux
-- C++ compiler : /usr/bin/c++
-- Release CXX flags: -O3 -DNDEBUG -fPIC -Wall -Wno-sign-compare -Wno-uninitialized
-- Debug CXX flags: -g -fPIC -Wall -Wno-sign-compare -Wno-uninitialized
-- Build type   : Release

-- BUILD_SHARED_LIBS : ON
-- BUILD_python     : ON
-- BUILD_matlab    : OFF
-- BUILD_docs      : ON
-- CPU_ONLY        : OFF
-- USE_OPENCV      : ON
-- USE_LEVELDB     : ON
-- USE_LMDB        : ON
-- USE_NCCL        : OFF
-- ALLOW_LMDB_NOLOCK : OFF
-- USE_HDF5        : ON
--
-- Dependencies:
-- BLAS          : Yes (Atlas)
-- Boost         : Yes (ver. 1.65)
-- glog          : Yes
-- gflags         : Yes
-- protobuf       : Yes (ver. 3.0.0)
-- lmdb          : Yes (ver. 0.9.21)
-- LevelDB       : Yes (ver. 1.20)
-- Snappy         : Yes (ver. ...)
-- OpenCV        : Yes (ver. 4.1.1)
-- CUDA          : Yes (ver. 10.0)
--
-- NVIDIA CUDA:
-- Target GPU(s)  : Auto
-- GPU arch(s)   : sm_72
-- cuDNN         : Yes (ver. 7.6.3)
--
-- Python:
-- Interpreter   : /usr/bin/python2.7 (ver. 2.7.17)
-- Libraries     : /usr/lib/aarch64-linux-gnu/libpython2.7.so (ver 2.7.17)
-- NumPy         : /usr/lib/python2.7/dist-packages/numpy/core/include (ver 1.13.3)
--
-- Documentation:
-- Doxygen       : No
-- config_file   :
--
-- Install:
-- Install path   : /home/nvidia/caffe/build/install
--
-- Configuring done
-- Generating done
-- Build files have been written to: /home/nvidia/caffe/build
[ 1%] Built target caffeprotobuf
Scanning dependencies of target caffe
[ 1%] Linking CXX shared library ../../lib/libcaffe.so
[ 98%] Built target caffe
[100%] Linking CXX shared library ../../lib/_caffe.so
Creating symlink /home/nvidia/caffe/python/caffe/_caffe.so -> /home/nvidia/caffe/build/lib/_caffe.so
[100%] Built target pycaffe
root@C:/home/nvidia/caffe/build#
```

Compile Caffe & Test

Jetson:

```
cd /media/green/M0/rootfs-xavier/home/nvidia/caffe  
sudo cp build/tools/ /mnt/rootfs/home/nvidia/caffe/build/ -rf  
sudo cp models examples data /mnt/rootfs/home/nvidia/caffe/ -rf  
sudo cp build/examples/ /mnt/rootfs/home/nvidia/caffe/build/  
  
cd caffe  
. ./data/mnist/get_mnist.sh  
. ./examples/mnist/create_mnist.sh  
  
. ./build/tools/caffe train -solver=examples/mnist/lenet_solver.prototxt
```

Compile Caffe & Test

```
nvidia@nvidia-desktop:~/caffe$ ./examples/mnist/create_mnist.sh
Creating lmdb...
I0420 10:02:35.221259 11284 db_lmdb.cpp:35] Opened lmdb examples/mnist/mnist_train_lmdb
I0420 10:02:35.221801 11284 convert_mnist_data.cpp:88] A total of 60000 items.
I0420 10:02:35.221853 11284 convert_mnist_data.cpp:89] Rows: 28 Cols: 28
I0420 10:02:36.328914 11284 convert_mnist_data.cpp:108] Processed 60000 files.
I0420 10:02:36.956005 11289 db_lmdb.cpp:35] Opened lmdb examples/mnist/mnist_test_lmdb
I0420 10:02:36.956511 11289 convert_mnist_data.cpp:88] A total of 10000 items.
I0420 10:02:36.956573 11289 convert_mnist_data.cpp:89] Rows: 28 Cols: 28
I0420 10:02:37.181118 11289 convert_mnist_data.cpp:108] Processed 10000 files.
Done.

nvidia@nvidia-desktop:~/caffe$ ./examples/mnist/train_lenet.sh
I0420 10:03:21.101969 11297 caffe.cpp:204] Using GPUs 0
I0420 10:03:21.132551 11297 caffe.cpp:209] GPU 0: Xavier
I0420 10:03:21.694084 11297 solver.cpp:45] Initializing solver from parameters:
test_iter: 100
test_interval: 500
base_lr: 0.01
display: 100
max_iter: 10000
lr_policy: "inv"
gamma: 0.0001
power: 0.75
momentum: 0.9
weight_decay: 0.0005
snapshot: 5000
snapshot_prefix: "examples/mnist/lenet"
solver_mode: GPU
device_id: 0
net: "examples/mnist/lenet_train_test.prototxt"
train_state {
    level: 0
    stage: ""
}
I0420 10:03:21.694917 11297 solver.cpp:102] Creating training net from net file: examples/mnist/lenet_train_test.prototxt
I0420 10:03:21.695561 11297 net.cpp:296] The NetState phase (0) differed from the phase (1) specified by a rule in layer mnist
I0420 10:03:21.695636 11297 net.cpp:296] The NetState phase (0) differed from the phase (1) specified by a rule in layer accuracy
I0420 10:03:21.695672 11297 net.cpp:53] Initializing net from parameters:
name: "LeNet"
state {
    phase: TRAIN
    level: 0
    stage: ""
}
layer {
    name: "mnist"
    type: "Data"
    top: "data"
    top: "label"
    include {
        phase: TRAIN
    }
    transform_param {
        scale: 0.00390625
    }
    data_param {
        source: "examples/mnist/mnist_train_lmdb"
        batch_size: 64
        backend: LMDB
    }
}
```

High load

The terminal window displays several pieces of information:

- System Load:** Shows CPU usage for 12 cores. Most cores are at 100% usage, with one core at 53.0% and another at 53.6%. The load average over 1, 5, and 15 minutes is 5.73, 2.23, and 3.82 respectively.
- Uptime:** The system has been running for 02:06:22.
- Memory Usage:** Total memory is 1.6GB, with 6G used, leaving 1.0GB free.
- Swap Usage:** Swap is currently not being used.
- Task List:** A detailed table shows 32 processes. Most are QEMU-related tasks (e.g., qemu-aarch64-static, cicc) running on various cores. Other processes include htop, Xorg, and various NVIDIA-related tools like nvcc, cuda, and caffe.
- Bottom Bar:** Contains navigation keys: F1Help, F2Setup, F3Search, F4Filter, F5Tree, F6SortBy, F7Nice, F8Nice, F9Kill, and F10Quit.

```

1 [|||] |99.3% 100.0%] 4 [|||] |100.0% 58.9%] 7 [|||] ||||| 7 [|||100.0%] 10 [|||] |100.0% 53.0%]
2 [|||] |100.0% 55.0%] 5 [|||] |100.0% 58.0%] 8 [|||] |100.0% 56.0%] 11 [|||] |100.0% 53.6%]
3 [|||] |100.0% 56.0%] 6 [|||] |100.0% 53.0%] 9 [|||] |100.0% 53.0%] 12 [|||100.0%] |100.0% 54.0%]
Mem[||| 1.6|||||6G] |||||s: 141, 242 thr; 2 running
Swp[]

Load average: 5.73 2.23 3.82
Uptime: 02:06:22

PID USER PRI NI VIRT RES SHR S CPU% MEM% TIME+ Command
1486 root 20 0 593M 263M 58452 12104 R 96.7 0.4 0:01.87 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1443 root 20 0 278M 75432 12104 R 94.7 0.5 0:02.85 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1431 root 20 0 279M 76432 12156 R 88.2 0.5 0:02.81 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1474 root 20 0 269M 64016 12120 R 85.5 0.4 0:02.40 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1404 root 20 0 306M 103M 12104 R 83.6 0.6 0:03.58 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1405 root 20 0 296M 93716 12104 R 82.2 0.6 0:03.33 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1419 root 20 0 292M 90556 11968 R 75.0 0.6 0:03.19 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1413 root 20 0 287M 85296 12236 R 71.7 0.5 0:03.12 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1464 root 20 0 271M 67112 12208 R 70.4 0.4 0:02.36 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1454 root 20 0 265M 59356 11984 R 69.7 0.4 0:02.07 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1467 root 20 0 264M 60256 12160 R 69.7 0.4 0:02.03 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1425 root 20 0 272M 67988 12160 R 66.4 0.4 0:02.32 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1480 root 20 0 263M 58620 12180 R 63.8 0.4 0:01.80 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1437 root 20 0 277M 73948 12176 R 61.2 0.5 0:02.62 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
1453 root 20 0 264M 59912 12200 R 60.5 0.4 0:02.05 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/..../nvvm/bin/cicc --c++14 --gnu_version=70400 --allow_managed --unsigned
31811 green 20 0 31696 5480 3448 R 0.7 0.0 0:00.38 htop
3373 green 20 0 341M 27240 21584 S 0.7 0.2 0:23.05 /usr/lib/x86_64-linux-gnu/xfce4/panel/plugins/libcpufreq.so 7 25165834
2265 root 20 0 297M 56144 39660 S 0.7 0.3 1:11.05 /usr/lib/xorg/Xorg -core :0 -seat seat0 -auth /var/run/lightdm/root/:0 -nolisten tcp vt7 -novtswitch
1471 root 20 0 217M 8612 4560 S 0.0 0.1 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1477 root 20 0 217M 7104 4560 S 0.0 0.0 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1483 root 20 0 217M 8888 4560 S 0.0 0.1 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1444 root 20 0 217M 7100 4560 S 0.0 0.0 0:00.03 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1398 root 20 0 217M 7100 4560 S 0.0 0.0 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1400 root 20 0 217M 7096 4560 S 0.0 0.0 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1410 root 20 0 217M 9124 4560 S 0.0 0.1 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1415 root 20 0 217M 8168 4560 S 0.0 0.0 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1422 root 20 0 217M 7108 4560 S 0.0 0.0 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1428 root 20 0 217M 7096 4560 S 0.0 0.0 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1434 root 20 0 217M 8340 4560 S 0.0 0.1 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1440 root 20 0 217M 8952 4560 S 0.0 0.1 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1447 root 20 0 217M 7096 4560 S 0.0 0.0 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1452 root 20 0 217M 7096 4560 S 0.0 0.0 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
1460 root 20 0 217M 7096 4560 S 0.0 0.0 0:00.02 /usr/bin/qemu-aarch64-static /bin/sh -c cicc --c++14 --gnu_version=70400 --allow_managed --unsigned_chars --diag_suppre
984 root 20 0 218M 10492 4928 S 0.0 0.1 0:00.06 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/nvcc /home/nvidia/caffe/src/caffe/layers/cudnn_relu_layer.cu -c -o /home
1138 root 20 0 218M 10176 4928 S 0.0 0.1 0:00.06 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/nvcc /home/nvidia/caffe/src/caffe/layers/cudnn_lcn_layer.cu -c -o /home
3381 green 20 0 678M 32688 25956 S 0.0 0.2 0:03.39 /usr/lib/x86_64-linux-gnu/xfce4/panel/plugins/libpulseaudio-plugin.so
1306 root 20 0 218M 9572 4928 S 0.0 0.1 0:00.06 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/nvcc /home/nvidia/caffe/src/caffe/layers/batch_norm_layer.cu -c -o /home
1327 root 20 0 218M 9360 4904 S 0.0 0.1 0:00.05 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/nvcc /home/nvidia/caffe/src/caffe/layers/cudnn_lrn_layer.cu -c -o /home
1238 root 20 0 218M 10968 4880 S 0.0 0.1 0:00.06 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/nvcc /home/nvidia/caffe/src/caffe/layers/cudnn_pooling_layer.cu -c -o /home
1252 root 20 0 218M 10484 4984 S 0.0 0.1 0:00.05 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/nvcc /home/nvidia/caffe/src/caffe/layers/bnll_layer.cu -c -o /home/nvid
1161 root 20 0 218M 9332 4932 S 0.0 0.1 0:00.06 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/nvcc /home/nvidia/caffe/src/caffe/layers/crop_layer.cu -c -o /home/nvid
1178 root 20 0 218M 10536 4928 S 0.0 0.1 0:00.05 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/nvcc /home/nvidia/caffe/src/caffe/layers/conv_layer.cu -c -o /home/nvid
1213 root 20 0 218M 9540 4880 S 0.0 0.1 0:00.05 /usr/bin/qemu-aarch64-static /usr/local/cuda/bin/nvcc /home/nvidia/caffe/src/caffe/layers/base_data_layer.cu -c -o /home
32413 root 20 0 231M 20500 10748 S 0.0 0.1 0:00.15 /usr/bin/qemu-aarch64-static /usr/local/bin/cmake -D verbose:BOOL= -D build_configuration:STRING=Release -D generated_fi
28338 green 20 0 123M 5932 4740 S 0.0 0.0 0:00.05 sshd: green@pts/10

```

Compile OpenCV & Test

```
-- PFM: YES
-- Video I/O:
--   DCI394: NO
--   FFmpeg: NO
--   avcodec: NO
--   avformat: NO
--   avutil: NO
--   swscale: NO
--   avresample: NO
--   GStreamer: YES (1.14.5)
--   v4l/v4l2: YES (linux/videodev2.h)
-- Parallel framework: pthreads
-- Trace: YES (with Intel ITT)
-- Other third-party libraries:
--   Lapack: NO
--   Eigen: YES (ver 3.3.4)
--   Custom HAL: YES (carotene (ver 0.0.1))
--   Protobuf: build (3.5.1)
-- NVIDIA CUDA: YES (ver 10.0, CUFFT CUBLAS)
--   NVIDIA GPU arch: 53 62 72
--   NVIDIA PTX archs:
-- cuDNN: YES (ver 7.6.3)
-- OpenCL:
--   Include path: /home/nvidia/opencv_clean/opencv/3rdparty/include/opencl/1.2
--   Link libraries: Dynamic load
-- Python 2:
--   Interpreter: /usr/bin/python2.7 (ver 2.7.17)
--   Libraries: /usr/lib/aarch64-linux-gnu/libpython2.7.so (ver 2.7.17)
--   numpy: /usr/lib/python2.7/dist-packages/numpy/core/include (ver 1.13.3)
--   install path: lib/python2.7/dist-packages/cv2/python-2.7
-- Python 3:
--   Interpreter: /usr/bin/python3 (ver 3.6.9)
--   Libraries: /usr/lib/aarch64-linux-gnu/libpython3.6m.so (ver 3.6.9)
--   numpy: /usr/local/lib/python3.6/dist-packages/numpy/core/include (ver 1.18.2)
--   install path: lib/python3.6/dist-packages/cv2/python-3.6
-- Python (for build): /usr/bin/python2.7
-- Java:
--   ant: NO
--   JNI: NO
--   Java wrappers: NO
--   Java tests: NO
-- Install to: /usr/local
-----
-- Configuring done
-- Generating done
-- Build files have been written to: /home/nvidia/opencv_clean/opencv/build
root@C:/home/nvidia/opencv_clean/opencv/build#
```

Compile OpenCV & Test

Compile OpenCV & Test

```
Scanning dependencies of target opencv_superres
[ 98%] Building CXX object modules/superres/CMakeFiles/opencv_superres.dir/src/super_resolution.cpp.o
[ 98%] Building CXX object modules/superres/CMakeFiles/opencv_superres.dir/src/btv_ll_cuda.cpp.o
[ 98%] Building CXX object modules/superres/CMakeFiles/opencv_superres.dir/src/frame_source.cpp.o
[ 98%] Building CXX object modules/superres/CMakeFiles/opencv_superres.dir/src/btv_ll.cpp.o
[ 98%] Building CXX object modules/superres/CMakeFiles/opencv_superres.dir/src/optical_flow.cpp.o
[ 98%] Building CXX object modules/superres/CMakeFiles/opencv_superres.dir/src/input_array_utility.cpp.o
[ 98%] Building CXX object modules/superres/CMakeFiles/opencv_superres.dir/opencl_kernels_superres.cpp.o
Scanning dependencies of target opencv_videostab
[ 98%] Building CXX object modules/videostab/CMakeFiles/opencv_videostab.dir/src/frame_source.cpp.o
[ 98%] Building CXX object modules/videostab/CMakeFiles/opencv_videostab.dir/src/fast_marching.cpp.o
[ 98%] Building CXX object modules/videostab/CMakeFiles/opencv_videostab.dir/src/inpainting.cpp.o
[ 98%] Building CXX object modules/videostab/CMakeFiles/opencv_videostab.dir/src/deblurring.cpp.o
[ 98%] Building CXX object modules/videostab/CMakeFiles/opencv_videostab.dir/src/motion_stabilizing.cpp.o
[ 98%] Building CXX object modules/videostab/CMakeFiles/opencv_videostab.dir/src/optical_flow.cpp.o
[ 98%] Building CXX object modules/videostab/CMakeFiles/opencv_videostab.dir/src/log.cpp.o
[ 98%] Building CXX object modules/videostab/CMakeFiles/opencv_videostab.dir/src/global_motion.cpp.o
[ 98%] Building CXX object modules/videostab/CMakeFiles/opencv_videostab.dir/src/stabilizer.cpp.o
[ 98%] Building CXX object modules/videostab/CMakeFiles/opencv_videostab.dir/src/outlier_rejection.cpp.o
[ 98%] Building CXX object modules/videostab/CMakeFiles/opencv_videostab.dir/src/wobble_suppression.cpp.o
[ 98%] Linking CXX shared library ../../lib/libopencv_superres.so
[ 98%] Built target opencv_superres
Scanning dependencies of target opencv_test_superres
Scanning dependencies of target opencv_perf_superres
[ 98%] Building CXX object modules/superres/CMakeFiles/opencv_test_superres.dir/test/test_main.cpp.o
[ 98%] Building CXX object modules/superres/CMakeFiles/opencv_test_superres.dir/test/test_superres.cpp.o
[ 98%] Building CXX object modules/superres/CMakeFiles/opencv_perf_superres.dir/perf/perf_main.cpp.o
[ 98%] Building CXX object modules/superres/CMakeFiles/opencv_perf_superres.dir/perf/perf_superres.cpp.o
[ 98%] Linking CXX executable ../../bin/opencv_test_superres
[ 98%] Built target opencv_test_superres
[ 99%] Linking CXX shared library ../../lib/libopencv_videostab.so
[ 99%] Built target opencv_videostab
Scanning dependencies of target opencv_test_videostab
Scanning dependencies of target opencv_python3
Scanning dependencies of target opencv_python2
[ 99%] Building CXX object modules/videostab/CMakeFiles/opencv_test_videostab.dir/test/test_stabilizer.cpp.o
[ 99%] Building CXX object modules/videostab/CMakeFiles/opencv_test_videostab.dir/test/test_main.cpp.o
[ 99%] Building CXX object modules/videostab/CMakeFiles/opencv_test_videostab.dir/test/test_motion_estimation.cpp.o
[ 99%] Building CXX object modules/python2/CMakeFiles/opencv_python2.dir/_/src2/cv2.cpp.o
[ 99%] Building CXX object modules/python3/CMakeFiles/opencv_python3.dir/_/src2/cv2.cpp.o
[ 99%] Linking CXX executable ../../bin/opencv_perf_superres
[ 99%] Built target opencv_perf_superres
[ 99%] Linking CXX executable ../../bin/opencv_test_videostab
[ 99%] Built target opencv_test_videostab
[ 99%] Linking CXX shared module ../../lib/python3/cv2.cpython-36m-aarch64-linux-gnu.so
[ 99%] Built target opencv_python3
[100%] Linking CXX shared module ../../lib/cv2.so
[100%] Built target opencv_python2
root@C:/home/nvidia/opencv_clean/opencv/build# `
```

Compile OpenCV & Test

```
/usr/local/bin/cmake -D WITH_CUDA=ON -D BUILD_CUDA_STUBS=ON -D OPENCV_EXTRA_MODULES_PATH=../../opencv_contrib/modules ..
```

Make -j16

```
sudo cp build/lib/cv2.so /mnt/rootfs/usr/local/lib/python2.7/dist-packages/
```

```
export LD_LIBRARY_PATH=/home/nvidia/opencv/lib:$LD_LIBRARY_PATH
```

Python

```
Import cv2
```

```
nvidia@nvidia-desktop:~/opencv/lib$ python
Python 2.7.17 (default, Nov  7 2019, 10:07:09)
[GCC 7.4.0] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> import cv2
>>> cv2.__version__
'4.3.0-dev'
>>> █
```

AGENDA

Jetson product and marketing

Why we need cross compile

How to setup the cross-compile environment

1. Config directly via the environment variable.
 2. Config in Makefile.
 3. Build your simulation.
-

Compile on X86 host and execute on Jetson

1. Cross compile Linux kernel.
2. Cross compile CUDA project.
3. Cross compile Multi Media API based application.
4. Cross compile Deepstream.
5. Cross compile Caffe
6. Cross compile OpenCV.

Where can I get the prebuild version and how to build on jetson?

JETSON DOWNLOAD CENTER

Prebuild Jetson Image:

<https://developer.nvidia.com/jetson-nano-sd-card-image>

```
Welcome to fdisk (util-linux 2.31.1).
Changes will remain in memory only, until you decide to write them.
Be careful before using the write command.
```

```
Command (m for help): p
Disk jetson-nano-sd-r32.1-2019-03-18.img: 12 GiB, 12884901888 bytes, 25165824 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: gpt
Disk identifier: D048AD43-24FD-4DED-B06E-7BB8ED98158C
```

Device	Start	End	Sectors	Size	Type
jetson-nano-sd-r32.1-2019-03-18.img1	24576	25165790	25141215	12G	Linux filesystem
jetson-nano-sd-r32.1-2019-03-18.img2	2048	2303	256	128K	Linux filesystem
jetson-nano-sd-r32.1-2019-03-18.img3	4096	4991	896	448K	Linux filesystem
jetson-nano-sd-r32.1-2019-03-18.img4	6144	7295	1152	576K	Linux filesystem
jetson-nano-sd-r32.1-2019-03-18.img5	8192	8319	128	64K	Linux filesystem
jetson-nano-sd-r32.1-2019-03-18.img6	10240	10623	384	192K	Linux filesystem
jetson-nano-sd-r32.1-2019-03-18.img7	12288	13439	1152	576K	Linux filesystem
jetson-nano-sd-r32.1-2019-03-18.img8	14336	14463	128	64K	Linux filesystem
jetson-nano-sd-r32.1-2019-03-18.img9	16384	17663	1280	640K	Linux filesystem
jetson-nano-sd-r32.1-2019-03-18.img10	18432	19327	896	448K	Linux filesystem
jetson-nano-sd-r32.1-2019-03-18.img11	20480	20735	256	128K	Linux filesystem
jetson-nano-sd-r32.1-2019-03-18.img12	22528	22687	160	80K	Linux filesystem

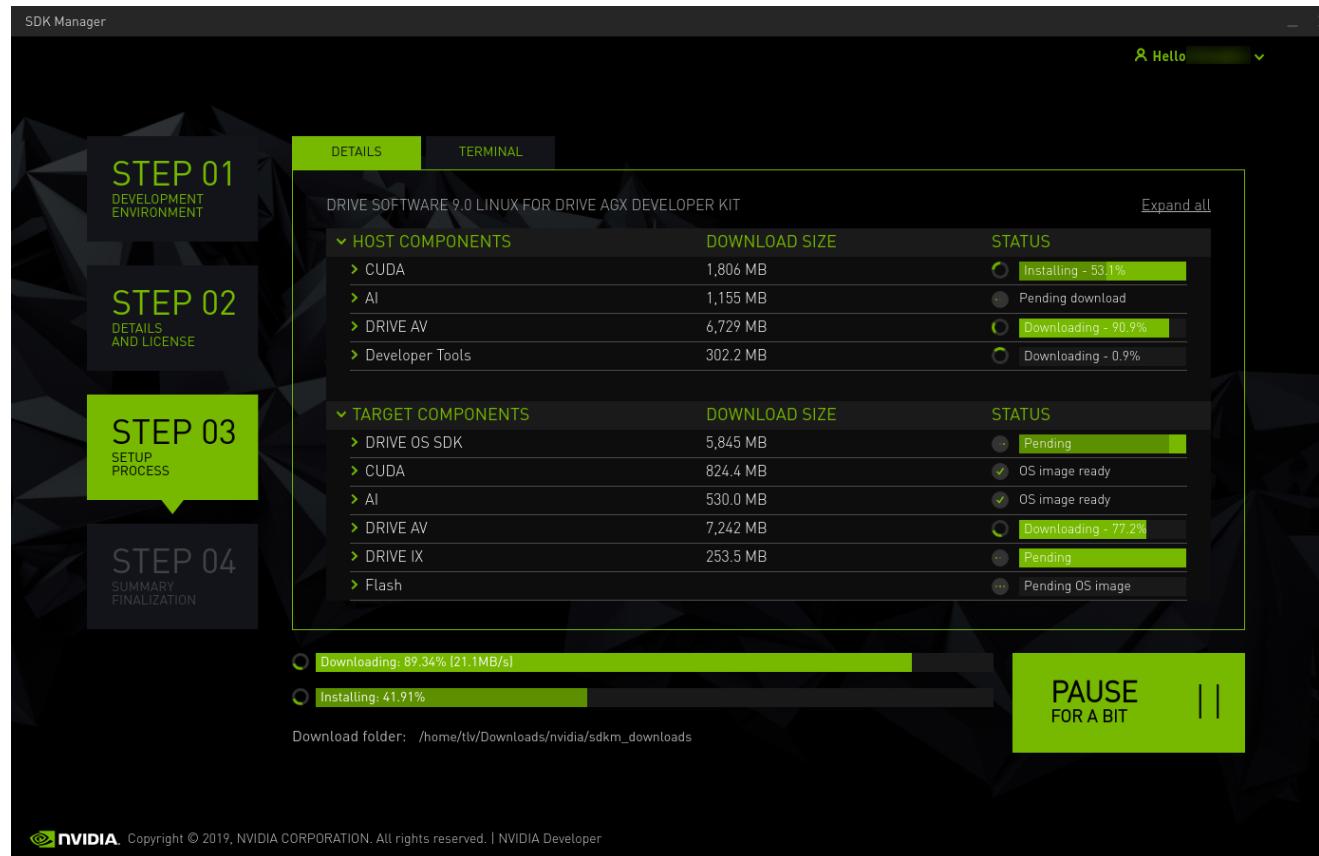
```
Partition table entries are not in disk order.
```

```
Command (m for help):
```

From here you can see the roots start from block 24576 with block size 512bytes. So We need to mount the image from offset
24576 * 512 = 12582912

```
root@C:~# mount -o loop,offset=12582912 jetson-nano-sd-r32.1-2019-03-18.img /mnt
root@C:~# ls /mnt/
```

SDKMANAGER



<https://developer.nvidia.com/nvidia-sdk-manager>

Backup and Store

To clone a Jetson device and flash

Clone:

- `sudo ./flash.sh -r -k APP -G <clone> <target> mmcblk0p1`

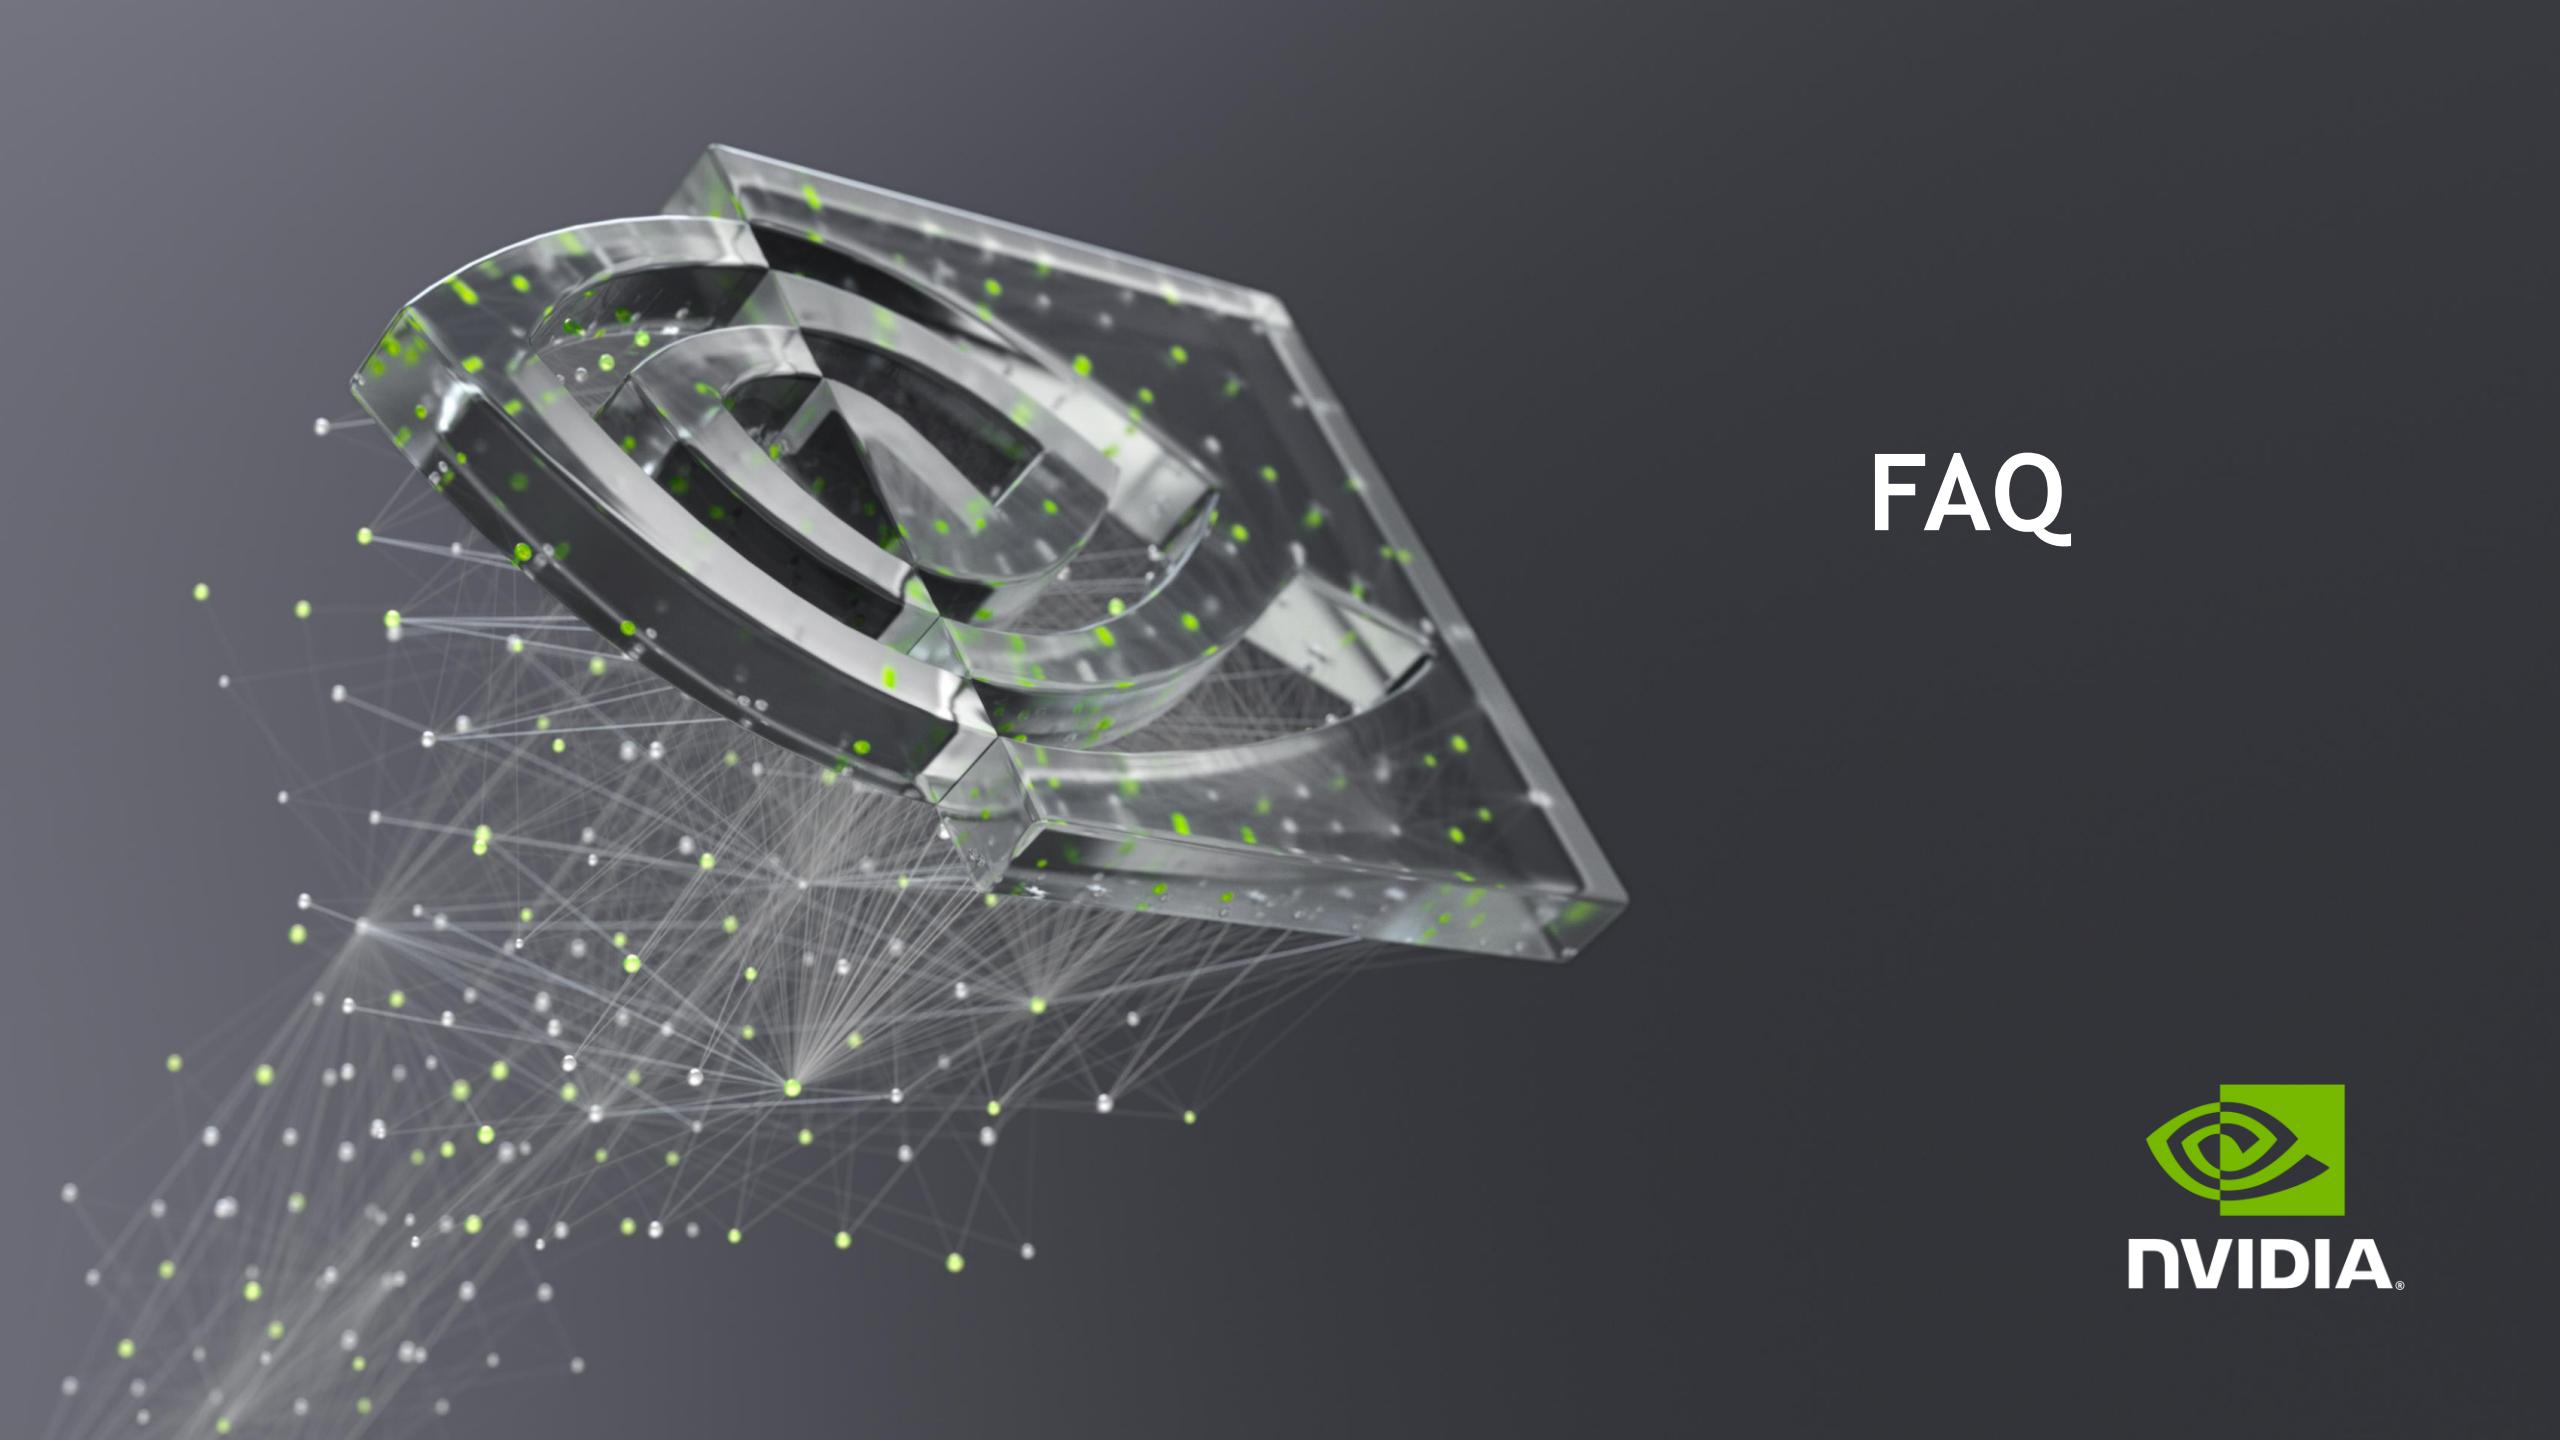
Customization:

- `simg2img system.img system.img.raw`
- `sudo mount -t ext4 -o loop system.img.raw /mnt`

Store:

- `sudo img2simg system.img.raw system.img`
- `sudo cp system.img bootloader/system.img`
- `sudo ./flash.sh -r -k APP <target> mmcblk0p1`
- Backup the ubuntu system and make a new Image based on it.

<https://docs.nvidia.com/jetson/l4t/index.html#page/Tegra%2520Linux%2520Driver%2520Package%2520Development%2520Guide%2Fflashing.html%23>



FAQ



QUESTIONS

1. How about chroot to a mounted nfs rootfs?
2. How about compiling tensorflow?